

Operating Systems

Lecture 6

Condition Variables

There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution.

Example:

A parent thread might wish to check whether a child thread has *completed*.

This is often called a `join()`.

A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

What we would like to see here is:

```
parent: begin
child
parent: end
```

Parent waiting for child: Spin-based Approach

```
1      volatile int done = 0;
2
3      void *child(void *arg) {
4          printf("child\n");
5          done = 1;
6          return NULL;
7      }
8
9      int main(int argc, char *argv[]) {
10         printf("parent: begin\n");
11         pthread_t c;
12         pthread_create(&c, NULL, child, NULL); // create child
13         while (done == 0)
14             ; // spin
15         printf("parent: end\n");
16         return 0;
17     }
```

This is hugely inefficient as the parent spins and **wastes CPU time**.

Condition variable

= Queue of threads

Waiting on the condition

An explicit queue that threads can put themselves on when some state of execution is not as desired.

Signaling on the condition

Some other thread, *when it changes its state*, can wake one of those waiting threads and allow them to continue.

Three in a package

condition variable `c`

state variable `m`

lock `L`; // to protect state variable

Declare condition variable

```
Pthread_cond_t c;
```

Proper initialization is required.

Operation (the POSIX calls)

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);    // wait()  
pthread_cond_signal(pthread_cond_t *c);                      // signal()
```

The wait() call takes a mutex as a parameter.

The wait() call release the lock and put the calling thread to sleep.

When the thread wakes up, it must re-acquire the lock.

Parent waiting for Child: Use a condition variable

```
1      int done = 0;
2      pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3      pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5      void thr_exit() {
6          pthread_mutex_lock(&m);
7          done = 1;
8          pthread_cond_signal(&c);
9          pthread_mutex_unlock(&m);
10     }
11
12     void *child(void *arg) {
13         printf("child\n");
14         thr_exit();
15         return NULL;
16     }
17
18     void thr_join() {
19         pthread_mutex_lock(&m);
20         while (done == 0)
21             pthread_cond_wait(&c, &m);
22         pthread_mutex_unlock(&m);
23     }
24
```


Parent waiting for Child: Use a condition variable

```
(cont.)  
25      int main(int argc, char *argv[]) {  
26          printf("parent: begin\n");  
27          pthread_t p;  
28          Pthread_create(&p, NULL, child, NULL);  
29          thr_join();  
30          printf("parent: end\n");  
31          return 0;  
32      }
```

Parent:

Creates the child thread and continues running itself.

Calls into `thr_join()` to wait for the child thread to complete.

Acquires the lock.

Checks if the child is done.

Puts itself to sleep by calling `wait()`.

Releases the lock.

Child:

Prints the message "child".

Calls `thr_exit()` to wake up the parent thread.

Grabs the lock.

Sets the state variable `done`.

Signals the parent thus waking it.

```
1      void thr_exit() {
2          pthread_mutex_lock(&m);
3          pthread_cond_signal(&c);
4          pthread_mutex_unlock(&m);
5      }
6
7      void thr_join() {
8          pthread_mutex_lock(&m);
9          pthread_cond_wait(&c, &m);
10         pthread_mutex_unlock(&m);
11     }
```

`thr_exit()` and `thr_join()` without variable `done`

Imagine the case where the *child runs immediately*.

The child will signal, but there is no thread asleep on the condition.

When the parent runs, it will call wait and be stuck.

No thread will ever wake it.

```
1      void thr_exit() {
2          done = 1;
3          Pthread_cond_signal(&c);
4      }
5
6      void thr_join() {
7          if (done == 0)
8              Pthread_cond_wait(&c);
9      }
```

The issue here is a subtle **race condition**.

The parent calls `thr_join()`.

The parent checks the value of `done`.

It will see that it is 0 and try to go to sleep.

Just before it calls `wait` to go to sleep, the parent is interrupted and the child runs.

The child changes the state variable `done` to 1 and signals.

But no thread is waiting and thus no thread is woken.

When the parent runs again, it sleeps forever.

Producer

Produce data items

Wish to place data items in a buffer

Consumer

Get data items out of the buffer to consume them in some way

Example: Multi-threaded web server

A producer puts HTTP requests to a work queue

Consumer threads take requests out of this queue and process them

The Put and Get Routines (Version 1)

```
1      int buffer;  
2      int count = 0;    // initially, empty  
3  
4      void put(int value) {  
5          assert(count == 0);  
6          count = 1;  
7          buffer = value;  
8      }  
9  
10     int get() {  
11         assert(count == 1);  
12         count = 0;  
13         return buffer;  
14     }
```

Only put data into the buffer when `count` is zero.

i.e., when the buffer is *empty*.

Only get data from the buffer when `count` is one.

i.e., when the buffer is *full*.

```
1      void *producer(void *arg) {
2          int i;
3          int loops = (int) arg;
4          for (i = 0; i < loops; i++) {
5              put(i);
6          }
7      }
8
9      void *consumer(void *arg) {
10         int i;
11         while (1) {
12             int tmp = get();
13             printf("%d\n", tmp);
14         }
15     }
```

Producer puts an integer into the shared buffer loops number of times.

Consumer gets the data out of that shared buffer.

A single condition variable `cond` and associated lock `mutex`

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1)                       // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                                // p4
11             Pthread_cond_signal(&cond);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
```



```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             if (count == 0)                        // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&cond);           // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

p1-p3: A producer waits for the buffer to be empty.

c1-c3: A consumer waits for the buffer to be full.


With just *a single producer* and *a single consumer*, the code works.


If we have **more than** one of producer and consumer?


③ C_1

③ C_2

③ P_1


buffer


ready Q





sleep queue: { }

③ C_1

③ C_2

③ P_1

C_1 executes first.




sleep queue: { }

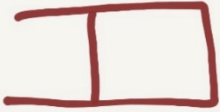
③ C_1


③ C_2

③ P_1

C_1 executes first.


buffer


ready Q






$\{C_1\}$

③ C_1

③ C_2

③ P_i

P_i executes. and send signal.






$\{C_1\}$

③ C_1

③ C_2

③ P_1

C_2 arrives.






$\{C_1\}$

③ C_1

③ C_2

③ P_1

C_1 gets signal.






~~{X}~~

③ C_1

③ C_2

③ P_1

C_2 consumes buffer.





{ }


③ C_1

③ C_2

③ P_1


buffer




 C_1
ready Q



{ P_1 }

③ C_1 ③ C_2 ③ P_1

C_1 reads the buffer!!!
empty

{ P_1 }

Reason: C_1 blindly reads the buffer
after it gets woken up.

The problem arises for a simple reason:

After the producer woke T_{c1} , but before T_{c1} ever ran, the state of the bounded buffer *changed by* T_{c2} .

There is no guarantee that when the woken thread runs, the state will still be as desired → Mesa semantics.

Virtually every system ever built employs *Mesa semantics*.

Hoare semantics provides a stronger guarantee that the woken thread will run immediately upon being woken.

Consumer T_{c1} wakes up and **re-checks** the state of the shared variable.

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == 1)                   // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                               // p4
11             Pthread_cond_signal(&cond);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
```

```
(Cont.)
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                    // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&cond);           // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

A simple rule to remember with condition variables is to **always use while loops**.


However, this code still has a bug (*next page*).


③ C_1

③ C_2

③ P_1


buffer


ready Q





sleep queue: { }

③ C_1

③ C_2

③ P_1

C_1 executes first.




sleep queue: { }

③ C_1


③ C_2

③ P_1

C_1 executes first.


buffer


ready Q




$\{C_1\}$

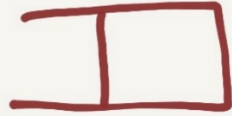
③ C_1


③ C_2

③ P_1

C_2 executes next.


buffer


ready Q





$\{C_1\}$

③ C_1

③ C_2

③ P_1

C_2 executes next





sleep queue : $\{C_1, C_2\}$

③ C_1

③ C_2

③ P_1

P_1 executes. and send signal.






Sleep queue: $\{C_1, C_2, P_1\}$

③ C_1

③ C_2

③ P_1

C_1 gets signal, executes and consumes buffer.





sleep queue : $\{P_1, C_2\}$

③ C_1

③ C_2

③ P_1

C_1 sends signal.




sleep queue : $\{P, C\}$

③ C_1


③ C_2

③ P_1

C_2 gets signal!


buffer


ready Q






sleep queue : $\{P_1, C_1\}$

③ C_1

③ C_2

③ P_1

C_2 gets executed. \rightarrow finds buffer empty. \rightarrow sleep.



sleep queue: $\{P_1, C_1, C_2\}$

Reason: C_2 gets signal.

P_1 should have gotten it!



Use **two** condition variables and while

Producer threads wait on the condition `empty`, and signals `fill`.

Consumer threads wait on `fill` and signal `empty`.

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&empty, &mutex);
10             put(i);
11             Pthread_cond_signal(&fill);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
15
```


The single Buffer Producer/Consumer Solution

```
(Cont.)  
16     void *consumer(void *arg) {  
17         int i;  
18         for (i = 0; i < loops; i++) {  
19             Pthread_mutex_lock(&mutex);  
20             while (count == 0)  
21                 Pthread_cond_wait(&fill, &mutex);  
22             int tmp = get();  
23             Pthread_cond_signal(&empty);  
24             Pthread_mutex_unlock(&mutex);  
25             printf("%d\n", tmp);  
26         }  
27     }
```

More **concurrency** and **efficiency** → Add more buffer slots.

Allow concurrent production or consuming to take place.

```
1      int buffer[MAX];
2      int fill = 0;
3      int use = 0;
4      int count = 0;
5
6      void put(int value) {
7          buffer[fill] = value;
8          fill = (fill + 1) % MAX;
9          count++;
10     }
11
12     int get() {
13         int tmp = buffer[use];
14         use = (use + 1) % MAX;
15         count--;
16         return tmp;
17     }
```

The Final Put and Get Routines

The Final Producer/Consumer Solution

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == MAX)                 // p2
9                  Pthread_cond_wait(&empty, &mutex); // p3
10             put(i);                               // p4
11             Pthread_cond_signal(&fill);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
```

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                    // c2
21                 Pthread_cond_wait(&fill, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&empty);          // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

The Final Working Solution (Cont.)

p2: **A producer** only sleeps if all buffers are currently filled.

c2: **A consumer** only sleeps if all buffers are currently empty.

The END