

Operating Systems

Lecture 9

13. The Abstraction: Address Space

Memory Virtualization

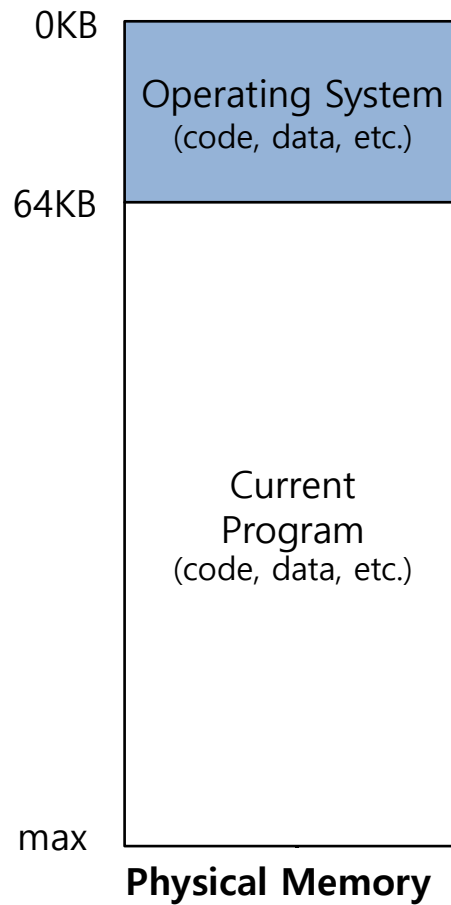
- ▣ What is **memory virtualization**?
 - ◆ OS virtualizes its physical memory.
 - ◆ OS provides an **illusion memory space** per each process.
 - ◆ It seems to be seen like **each process uses the whole memory**.

Benefit of Memory Virtualization

- ▣ Ease of use in programming
- ▣ Memory efficiency in terms of **times** and **space**
- ▣ The guarantee of isolation for processes as well as OS
 - ◆ Protection from **errant accesses** of other processes

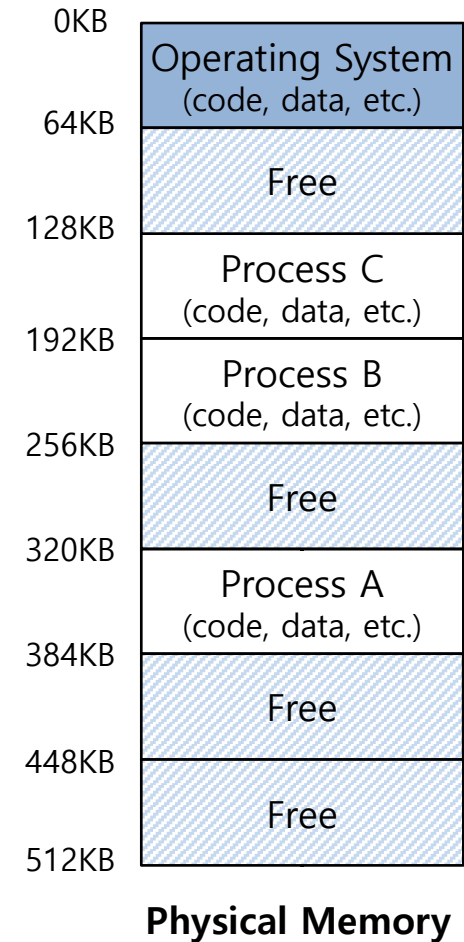
OS in The Early System

- ▣ Load only one process in memory.
 - ◆ Poor utilization and efficiency




Multiprogramming and Time Sharing

- ❑ **Load multiple processes** in memory.
 - ◆ Execute one for a short while.
 - ◆ Switch processes between them in memory.
 - ◆ Increase utilization and efficiency.
- ❑ Cause an important **protection issue**.
 - ◆ Errant memory accesses from other processes




Virtual Address Space

- OS creates an **abstraction** of physical memory.
 - ◆ The address space contains all about a running process.
 - ◆ That consists of program code, heap, stack and etc.



Virtual Address Space(Cont.)

- ▣ Code
 - ◆ Where instructions live
- ▣ Heap
 - ◆ Dynamically allocate memory.
 - `malloc` in C language
 - `new` in object-oriented language
- ▣ Stack
 - ◆ Store return addresses or values.
 - ◆ Contain local variables arguments to routines.



Virtual Address Space(Cont.)

▣ Process memory layout

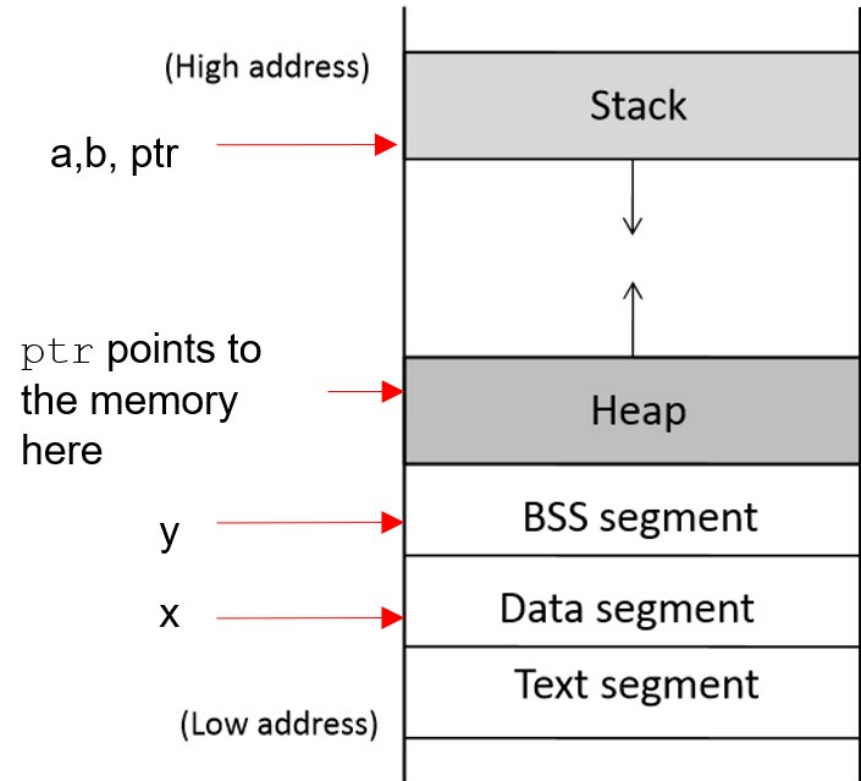
```
int x = 100;
int main()
{
    // data stored on stack
    int a=2;
    float b=2.5;
    static int y;

    // allocate memory on heap
    int *ptr = (int *) malloc(2*sizeof(int));

    // values 5 and 6 stored on heap
    ptr[0]=5;
    ptr[1]=6;

    // deallocate memory on heap
    free(ptr);

    return 1;
}
```



- ▣ **Every address** in a running program is virtual.

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code   : %p\n", (void *) main);
    printf("location of heap   : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack  : %p\n", (void *) &x);

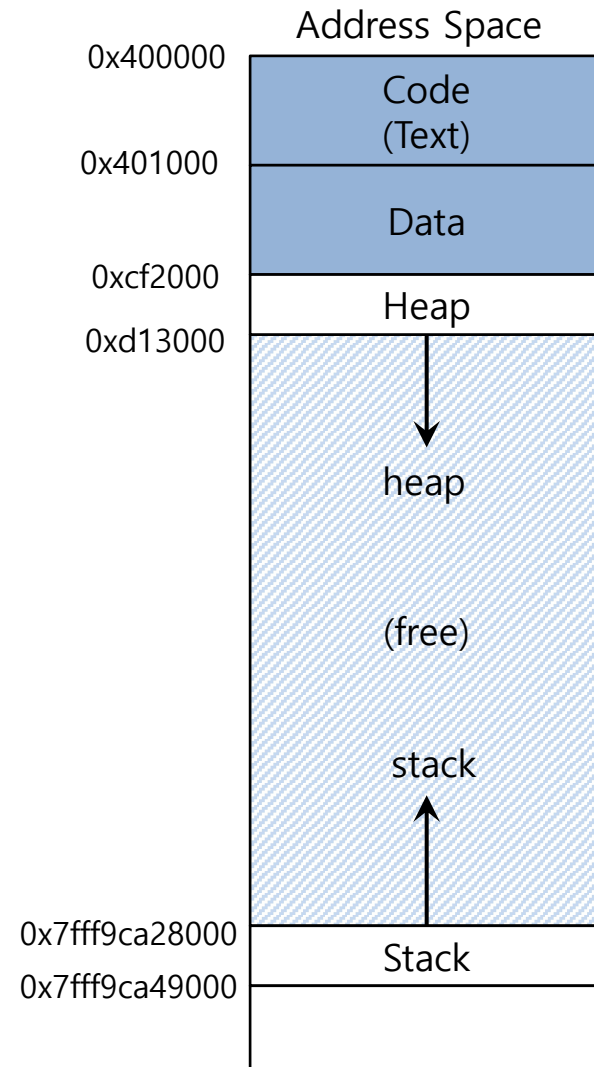
    return x;
}
```

A simple program that prints out addresses

Virtual Address(Cont.)

■ The output in 64-bit Linux machine

```
location of code   : 0x40057d  
location of heap   : 0xcf2010  
location of stack  : 0x7fff9ca45fcc
```

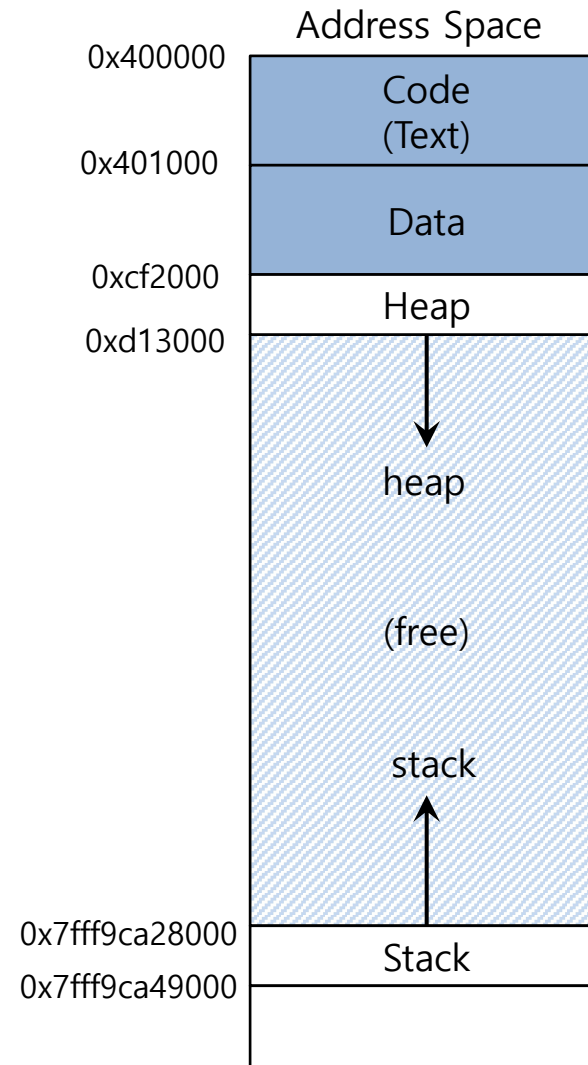


Virtual Address(Cont.)

■ The output in 64-bit Linux machine

```
location of code   : 0x40057d
location of heap   : 0xcf2010
location of stack  : 0x7fff9ca45fcc
```

OS translates the virtual address to physical address



15. Address Translation

Address Translation

- ▣ Hardware transforms a **virtual address** to a **physical address**.
 - ◆ The desired information is actually located in a physical address.
- ▣ The OS must get involved at key points to set up the hardware.
 - ◆ The OS must manage memory to judiciously intervene.

Example: Address Translation

▣ C - Language code

```
void func() {  
    int x=3000;  
    ...  
    x = x + 3; // this is the line of code we are interested in  
}
```

- ◆ **Load** a value from memory
- ◆ **Increment** it by three
- ◆ **Store** the value back into memory


Example: Address Translation(Cont.)

▣ Assembly

```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax
132 : addl $0x03, %eax         ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)      ; store eax back to mem
```

- ◆ **Load** the value at that address into `eax` register.
- ◆ **Add** 3 to `eax` register.
- ◆ **Store** the value in `eax` back into memory.

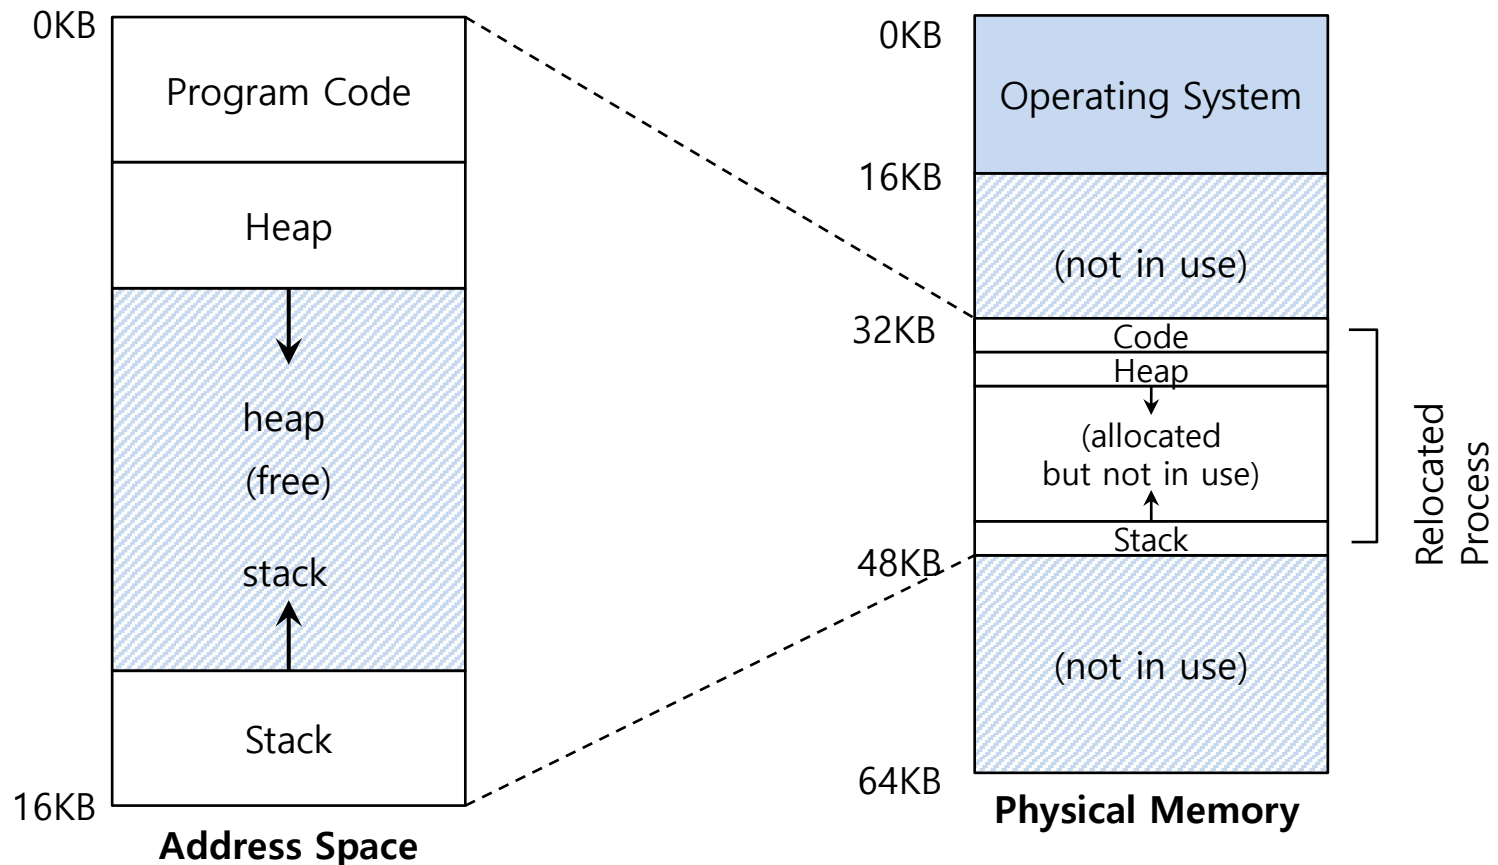
Example: Address Translation(Cont.)



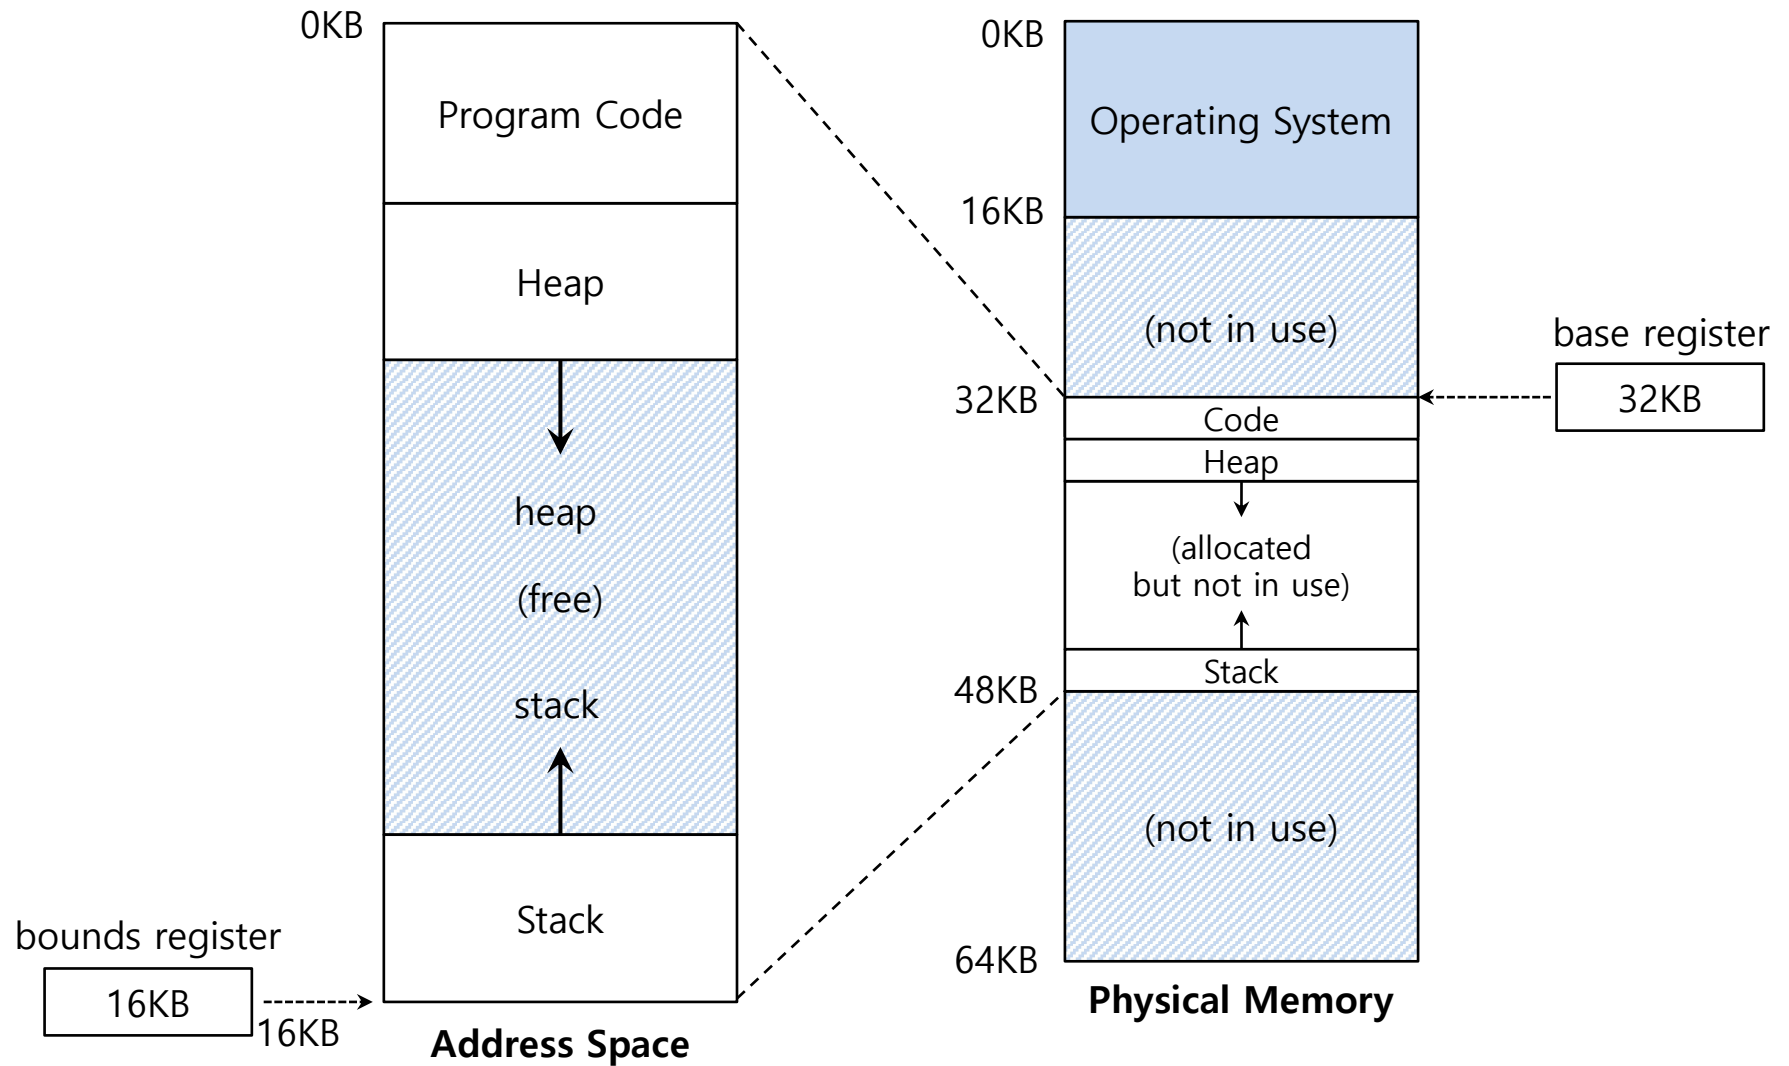
- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

Dynamic Relocation: Base and Bound Register

- The OS wants to place the process **somewhere else** in physical memory, not at address 0.
 - ◆ The address space start at address 0.



Base and Bounds Register



Dynamic(Hardware base) Relocation

- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**.
 - ◆ Set the **base** register a value.

$$\text{physical address} = \text{virtual address} + \text{base}$$

- ◆ Every virtual address must **not be greater than bound** and **negative**.

$$0 \leq \text{virtual address} < \text{bounds}$$

Relocation and Address Translation

128 : movl 0x0(%ebx), %eax


- ◆ **Fetch** instruction at address 128

$$32896 = 128 + 32KB(base)$$

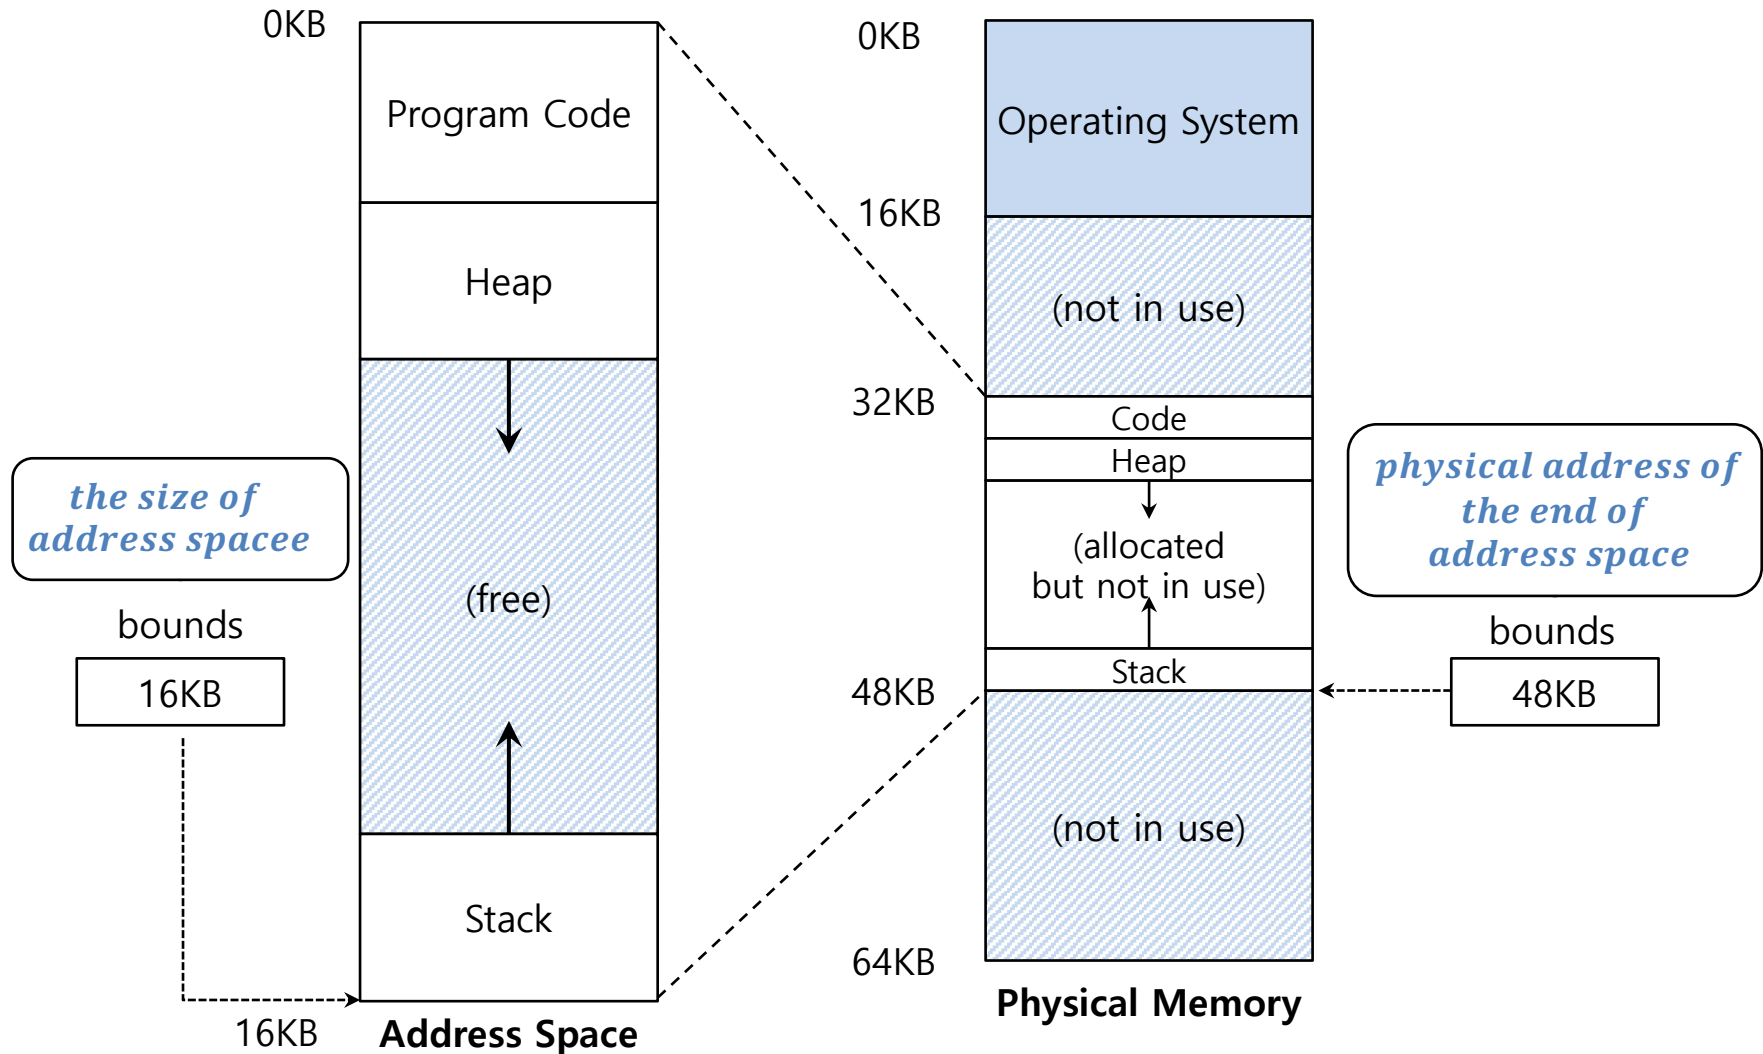
- ◆ **Execute** this instruction

- Load from address 15KB

$$47KB = 15KB + 32KB(base)$$



Bounds Register



Hardware Requirements

Hardware Requirements	Notes
Privileged mode	<i>Needed to prevent user-mode processes from executing privileged operations</i>
Base/bounds registers	<i>Need pair of registers per CPU to support address translation and bounds checks</i>
Ability to translate virtual addresses and check if within bounds	<i>Circuitry to do translations and check limits; in this case, quite simple</i>
Privileged instruction(s) to update base/bounds	<i>OS must be able to set these values before letting a user program run</i>
Privileged instruction(s) to register exception handlers	<i>OS must be able to tell hardware what code to run if exception occurs</i>
Ability to raise exceptions	<i>When processes try to access privileged instructions or out-of-bounds memory</i>

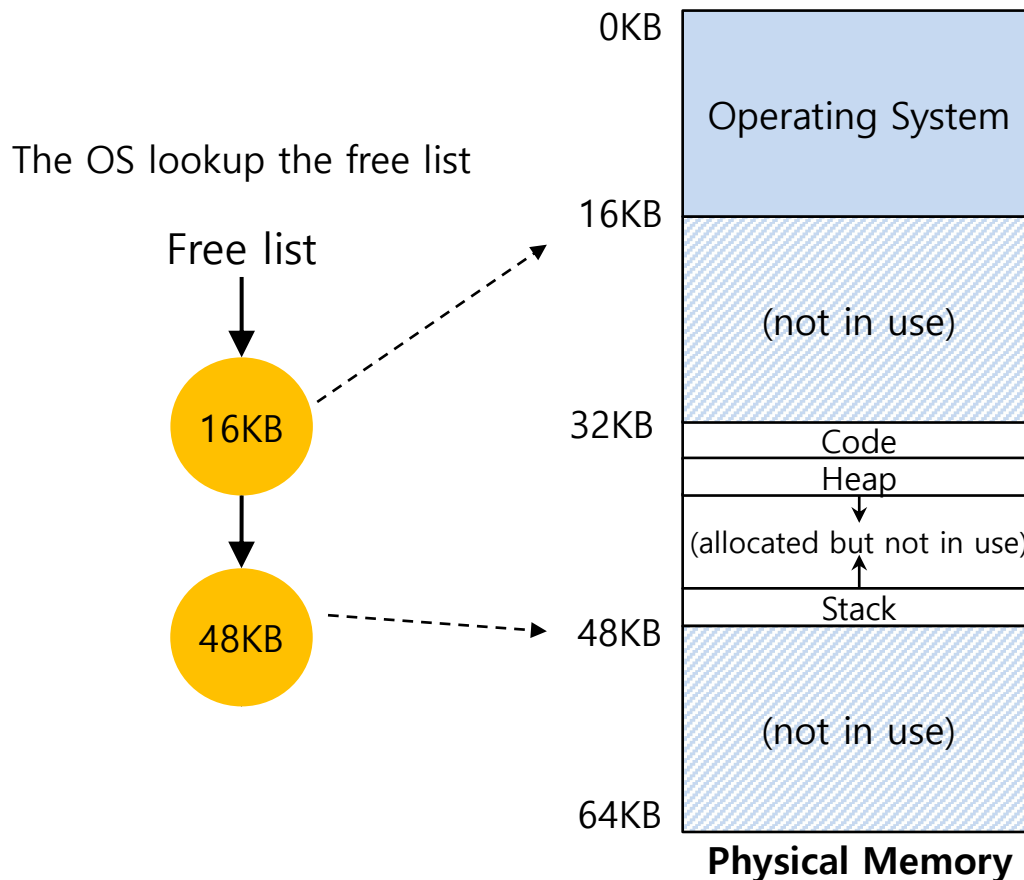
Dynamic Relocation: Hardware Requirements

OS Issues for Memory Virtualizing

- ▣ The OS must **take action** to implement **base-and-bounds** approach.
- ▣ Three critical actions:
 - ◆ When a process **starts running**:
 - Finding space for address space in physical memory
 - ◆ When a process is **terminated**:
 - Reclaiming the memory for use
 - ◆ When context **switch occurs**:
 - Saving and storing the base-and-bounds pair



OS Issues: When a Process Starts Running

- ▣ The OS must **find a room** for a new address space.
 - ◆ free list : A list of the range of the physical memory which are not in use.



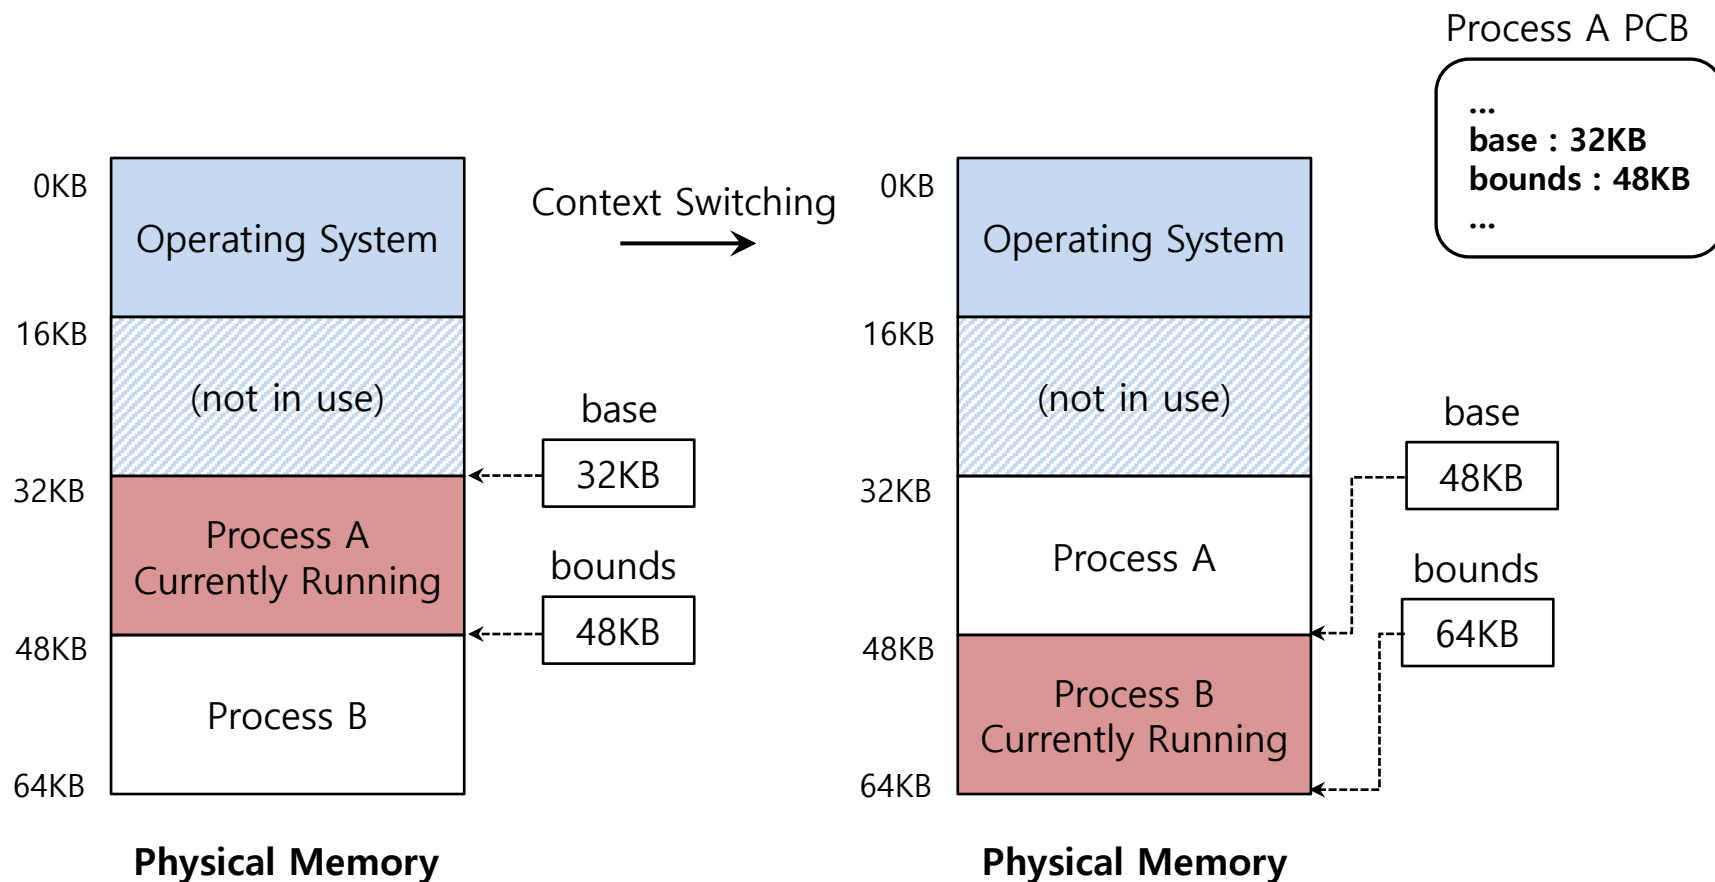
OS Issues: When a Process Is Terminated

- ▣ The OS must **put the memory back** on the free list.



OS Issues: When Context Switch Occurs

- The OS must **save and restore** the base-and-bounds pair.
 - ◆ In **process structure** or **process control block(PCB)**



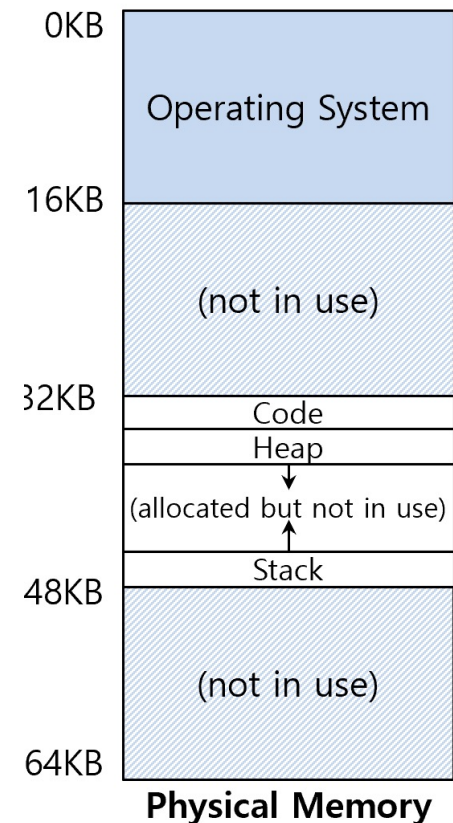
OS Issues: provide exception handlers

- ▣ the OS must provide exception handlers,
- ▣ the OS installs these handlers at boot time (via privileged instructions
 - ◆ Exception handler for segmentation fault

Inefficiency of Base and Bound registers

▣ Internal fragmentation

- ◆ The relocated process is using physical memory from 32 KB to 48 KB
- ◆ The process stack and heap are not too big, all of the space between the two is simply *wasted*.




Summary

- ▣ Address translation: hardware support and OS support
- ▣ Basic form: base and bound
- ▣ Fragmentation issue

16. Segmentation

Inefficiency of the Base and Bound Approach




- ▣ **Big chunk of “free” space**
- ▣ “free” space **takes up** physical memory.
- ▣ Hard to run when an address space **does not fit** into physical memory

Segmentation

- ▣ Segment is just a **contiguous portion** of the address space of a particular length.
 - ◆ Logically-different segment: code, stack, heap
- ▣ Each segment can be **placed** in **different part of physical memory**.
 - ◆ **Base** and **bounds** exist **per each segment**.

Placing Segment In Physical Memory




Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Address Translation on Segmentation: code

$$\text{physical address} = \text{offset} + \text{base}$$

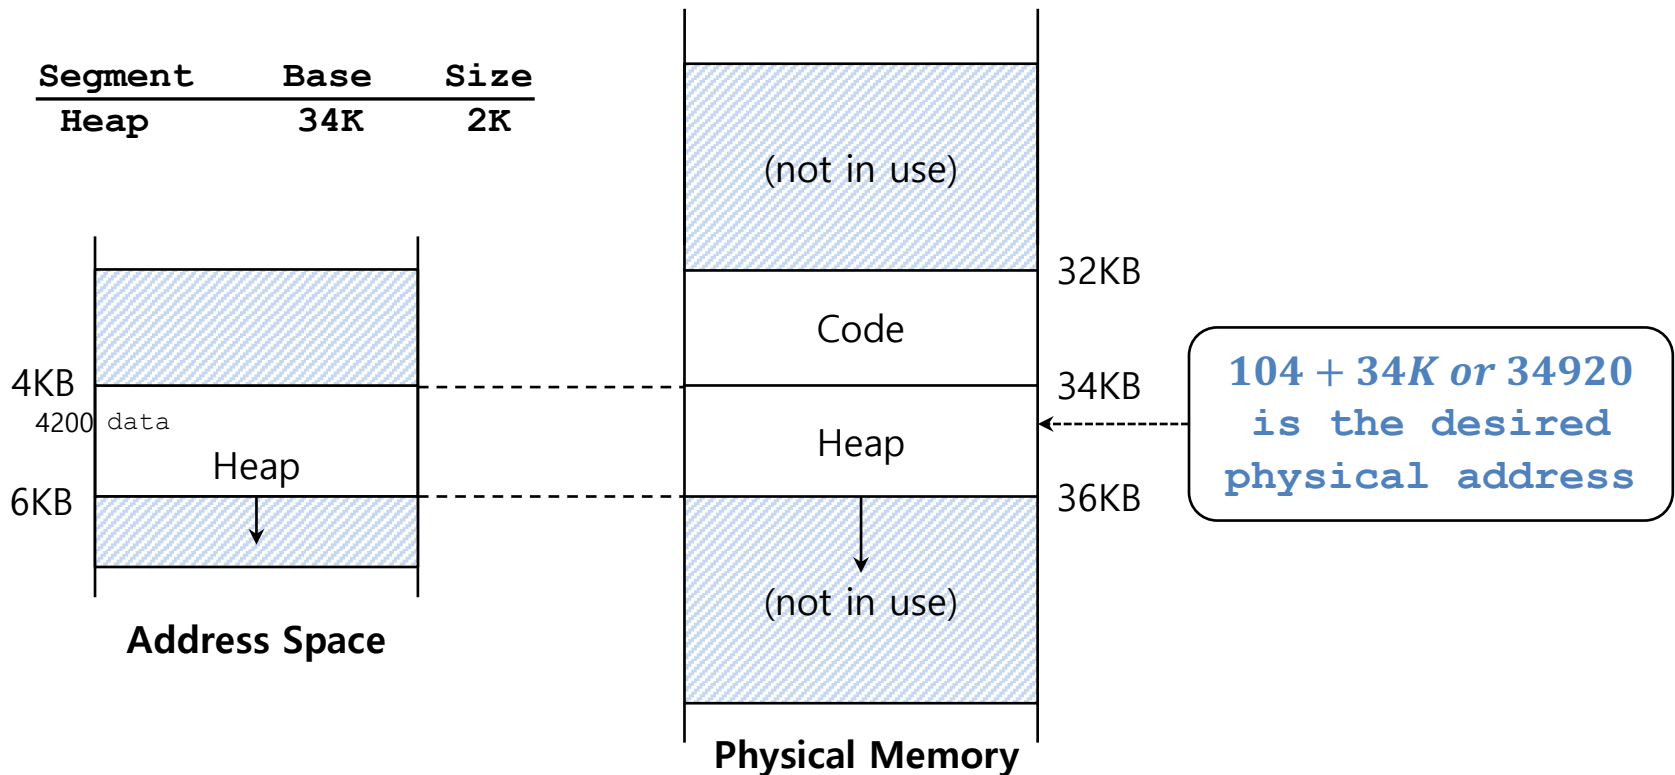
- ▣ The `offset` of virtual address 100 is 100.
 - ◆ The code segment **starts at virtual address 0** in address space.



Address Translation on Segmentation: heap


*Virtual address + base is not the correct physical address.
OFFSET of Virtual address + base is the correct physical address.*

- ▣ The offset of virtual address 4200 is 104.
 - ◆ The heap segment **starts at virtual address 4096** in address space.



Segmentation Fault or Violation


- If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS occurs **segmentation fault**.
 - ◆ The hardware detects that address is **out of bounds**.



Referring to Segment


□ Explicit approach

- ◆ Chop up the address space into segments based on the **top few bits** of virtual address.



□ Example: virtual address 4200 (010000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Segment selection

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- ◆ `SEG_MASK = 0x3000 (11000000000000)`
- ◆ `SEG_SHIFT = 12`
- ◆ `OFFSET_MASK = 0xFFF (00111111111111)`

Referring to Stack Segment

- ❑ Stack grows **backward**.
- ❑ **Extra hardware support** is need.
 - ◆ The hardware checks which way the segment grows.
 - ◆ 1: positive direction, 0: negative direction



Segment Register(with Negative-Growth Support)

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Support for Sharing

- Segment can be **shared between address** space.
 - ◆ **Code sharing** is still in use in systems today.
 - ◆ by extra hardware support.
- Extra hardware support is need for form of **Protection bits**.
 - ◆ **A few more bits** per segment to indicate **permissions** of **read**, write and **execute**.

Segment Register Values(with Protection)

Segment	Base	Size	Grows Positive?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write



Fine-Grained and Coarse-Grained segmentation

- ▣ **Coarse-Grained** means small number of segments.
 - ◆ e.g., code, heap, stack.
- ▣ **Fine-Grained** segmentation allows **more flexibility** for address space in some early system.
 - ◆ To support many segments, Hardware support with a **segment table** is required.

- ▣ **External Fragmentation:** little holes of **free space** in physical memory that is too small for allocating segment.
 - ◆ There is **24KB free**, but **not in one contiguous** segment.
 - ◆ The OS **cannot** satisfy the **20KB request**.

- ▣ **Compaction: rearranging** the exiting segments in physical memory.
 - ◆ Compaction is **costly**.
 - **Stop** running process.
 - **Copy** data to somewhere.
 - **Change** segment register value.

Memory Compaction



History of segmentation

- ▣ In early days, OS used segmentation.
 - ◆ Burroughs B5000 (first commercial machine with virtual memory)
 - ◆ IBM AS/400
 - ◆ Intel 8086, 80286
- ▣ 80386 and later Intel CPU's support paging.
- ▣ x86-64 does not use segmentation any more in 64bit mode
 - ◆ CS,SS,DS and ES are forced to 0 and 2^{24} ..

Summary

- ▣ Segmentation can better support sparse address spaces.
- ▣ It is also fast as the overheads of translation are minimal.
- ▣ Sharing (such as code) is easy.
- ▣ Issues
 - ◆ External fragmentation issue
 - ◆ Sparse segment