Operating Systems

Lecture 15

43. Log-structured File Systems

Overview

- In the early 90`s, a new file system known as the log-structured file system(LFS) was developed.
- The Motivation …
  - Memory sizes were growing.
  - Large gap between random IO and sequential IO performance.
  - Existing file system perform poorly on common workloads.
  - File systems were not RAID-aware. There exists small write problem.
- In this chapter, we study Log-Structured File System(LFS).
  - How can a file system **transform all writes into sequential writes**?

Writing to Disk Sequentially

- How do we transform all updates to file-system state into a series of sequential writes to disk?
  - data update

  - metadata needs to be updated too. (Ex. inode)

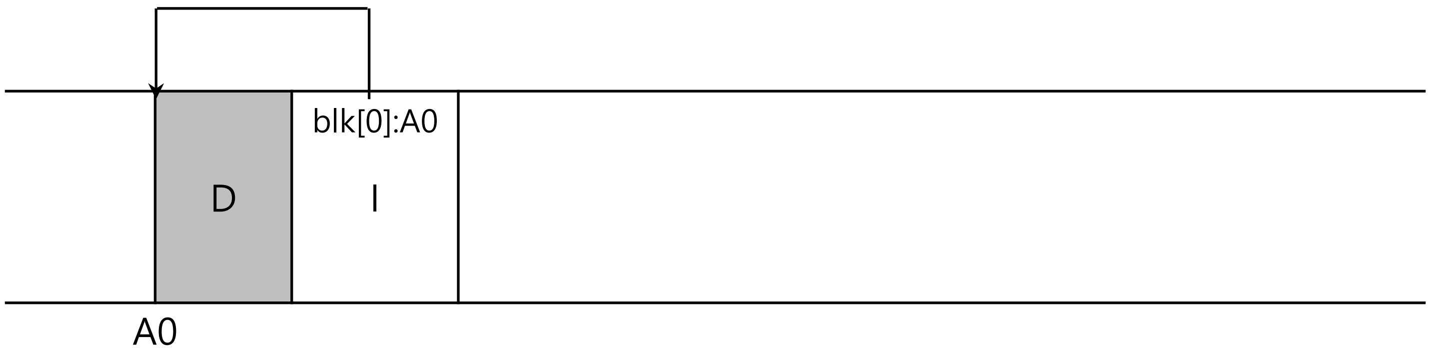

Segment

- Writing to the disk sequentially is not enough to guarantee the efficient writes.
  - Disk may rotate between the writes. → loose a single revolution between the writes.
- Write buffering.
  - Segment: a set of sequential writes that are written to the disk with a single unit.
  - Keep track of updates in **memory buffer**. (a few Mbyte)
  - Write them to disk all at once, when it has sufficient number of updates.

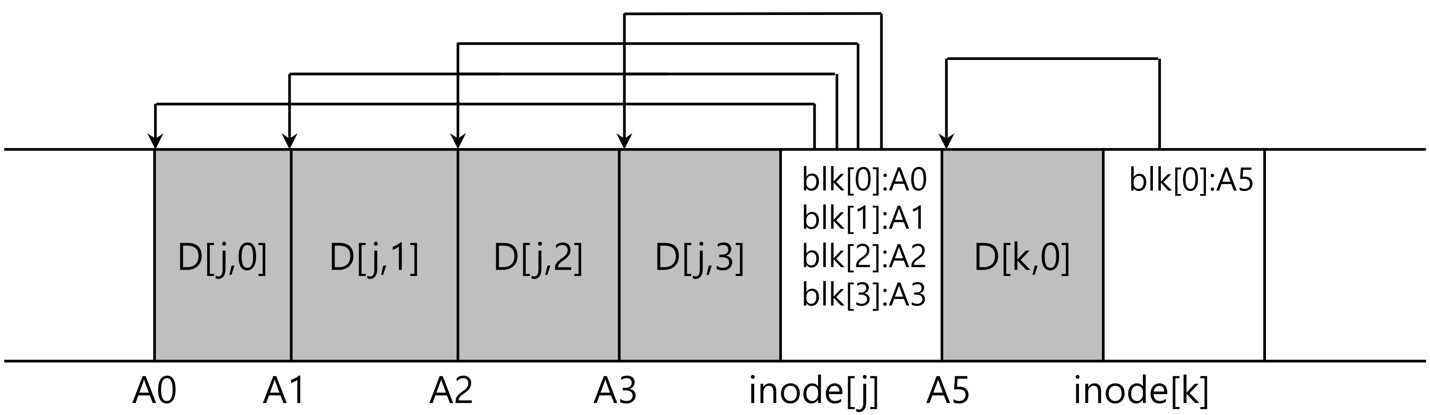

Right segment size

- Time to write D Mbyte

$$T_{write} = T_{position} + \frac{D}{R_{peak}}$$

- Effective write bandwidth

$$R_{effecitve} = \frac{D}{T_{write}} = \frac{D}{T_{position+} \frac{D}{R_{peak}}}$$

- We like to make the effective write bandwidth close to peak bandwidth with some fraction F (0<F<1)

$$R_{effecitve} = \frac{D}{T_{position+} \frac{D}{R_{peak}}} = \mathrm{F\ x}\ R_{peak}$$

Right segment size

- Then, D can be computed as follows.

$$D = \mathrm{F} \text{ x } R_{peak} \text{ x } (T_{position} + \frac{D}{R_{peak}})$$

$$D = \left(\mathrm{F} \text{ x } R_{peak} \text{ x } T_{position}\right) + \left(F \text{ x } R_{peak} \text{ x } \frac{D}{R_{peak}}\right)$$

$$D = \frac{F}{1-F} \text{ x } R_{peak} \text{ x } T_{position}$$

- Example: Positioning time 10 msec, peak transfer rate 100MByte/sec, we like to achieve 90% of the peak rate

D = 0.9*0.1*100 Mbyte/sec * 0.01 secs = 9 Mbyte

  - What is D if F = 0.95?

Finding Inode

- The position of the inodes keep changing.
- The Inode Map
    - A data structure that contains the location of the most recent inode for a given inode number.
    - Places the chunk of updated inode map next to the updated inode.
    - Where to find the inode map?

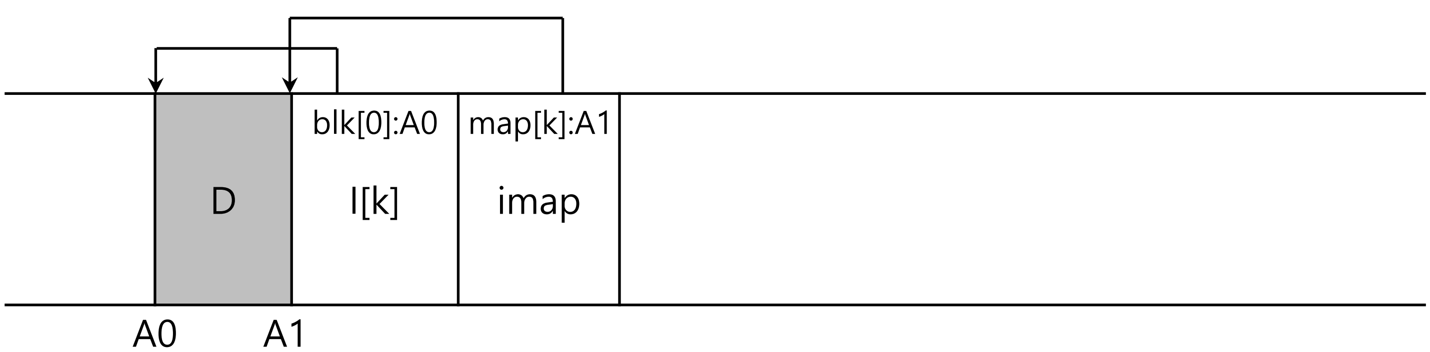

The Checkpoint Region

- How to find the inode map spread across the disk?
  - The LFS File system must have fixed location on disk to begin a file lookup.
- **Checkpoint Region**
  - fixed location in the LFS partition.
  - Contain the pointers to the latest of the inode map.

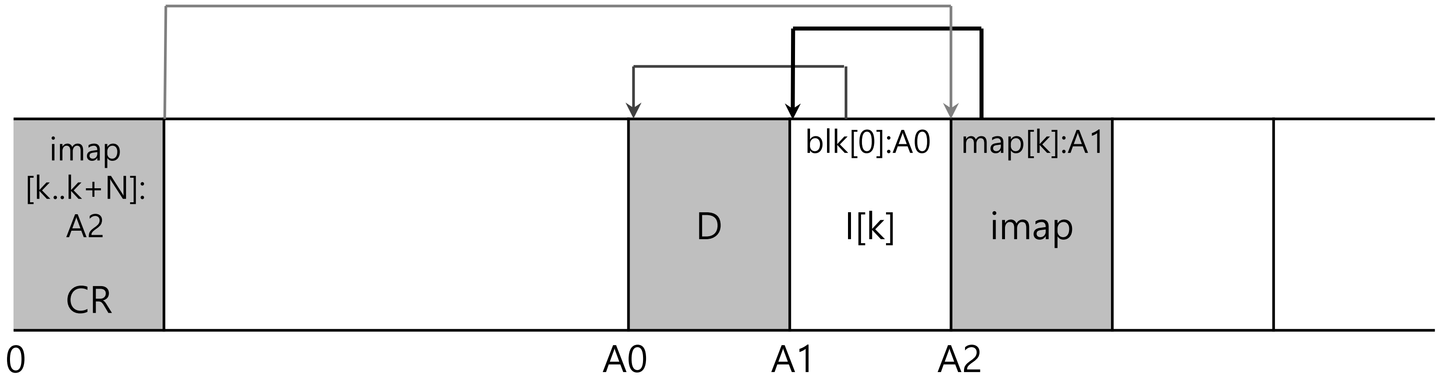

Reading a file from the disk

- Reading a file block
  - Read a checkpoint region
  - Read inode map
  - Read inode
  - Read data block
- What about sequential read?
  - It may become random read.

LFS is optimized for the write operation.

Garbage

- LFS keeps writing newer version of file.
- Garbage: LFS leaves the older versions of file structures all over the disk.
- An example of garbage
  - Overwrite the data block:

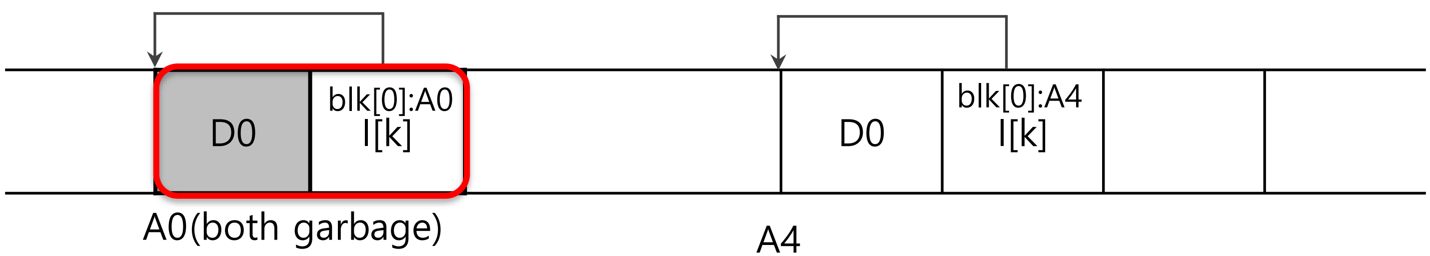

  - Append a block to that original file k:

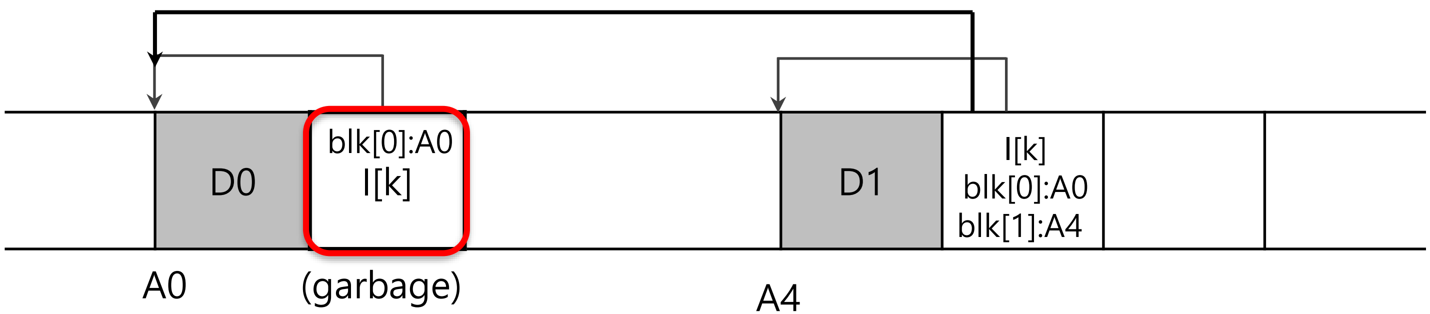

Garbage Collection (segment cleaning)

- What to do with the older versions of the block
  - Versioning filesystem: keep the old blocks and allow the users to restore to the older version of the filesystem status.
  - LFS: periodically clean the older versions of the file data, inodes and other structures.
- Unit of garbage collection: Segment
  - Reads a number of old segments (M segments).
  - Identify the valid blocks.
  - Write them to a number of new segments (in memory), N segments.
  - Write N segments to the disk.
  - Then, $N < M$.

Garbage collection

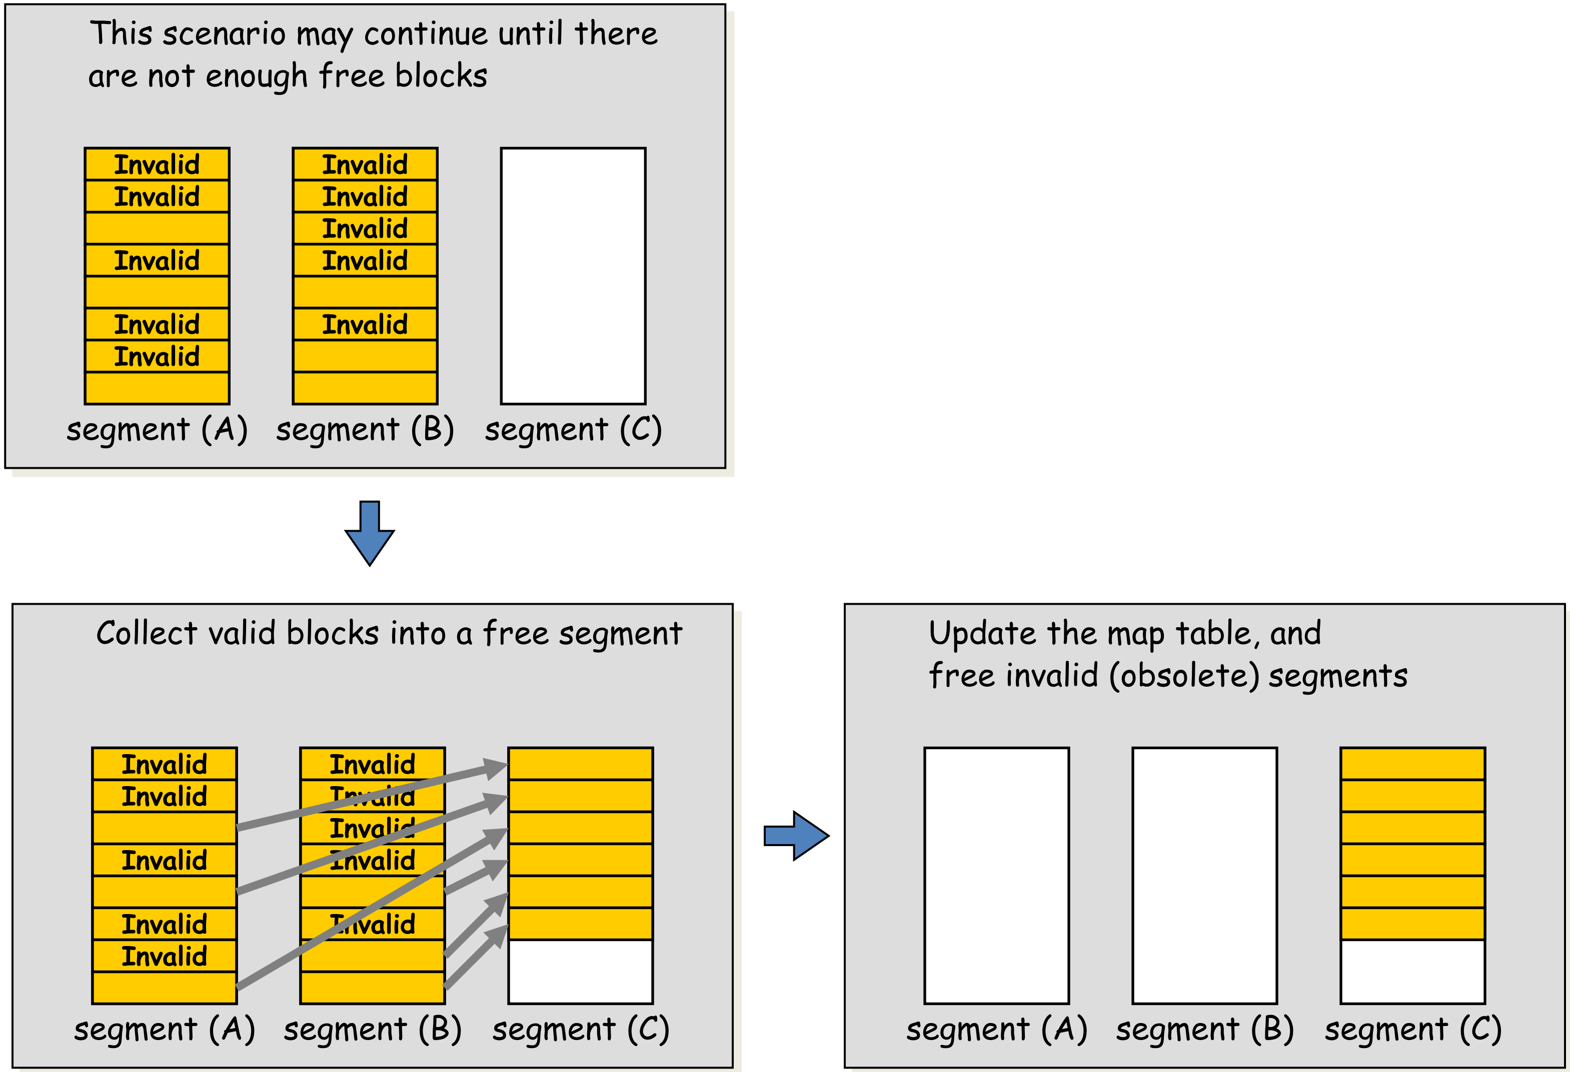

Summary

- Gather all updates into an in-memory segment.
  - Write them out together sequentially.
- LFS-style is excellent for performance on many different devices.
  - Hard drives, parity-based RAIDs, even Flash-based SSDs.
- Some modern commercial filesystems adopt a similar approach even though it generates garbage.
  - NetApp`s **WAFL,** Sun`s **ZFS** and Linux **btrfs**
  - In particular, WAFL turns cleaning problem into a feature, by providing old version of the file system via **snapshot.**

Thanks