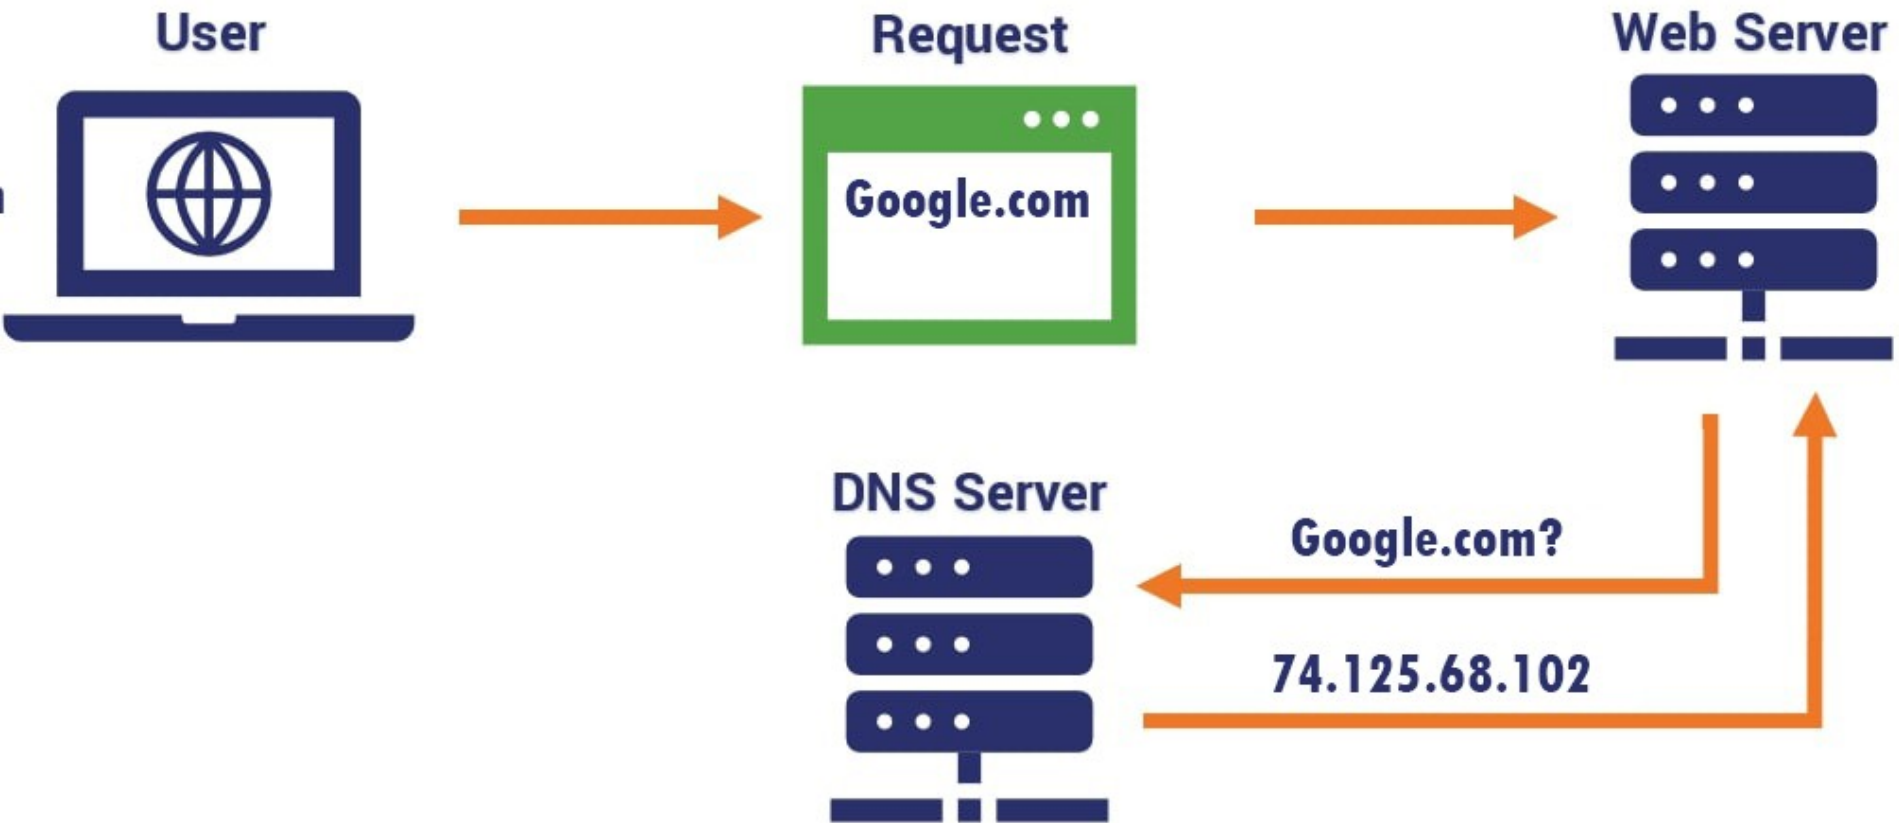


# Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP






# Function

**m** map between (domain name, service) to value, e.g.,

- (www.cs.yale.edu, Addr)  
→ 128.36.229.30
- (cs.yale.edu, Email)  
→ netra.cs.yale.edu



```
Reply from 103.235.46.39: bytes=32 time=3ms TTL=56
Reply from 103.235.46.39: bytes=32 time=3ms TTL=56
Reply from 103.235.46.39: bytes=32 time=3ms TTL=56
Reply from 103.235.46.39: bytes=32 time=3ms TTL=56
```

Ping statistics for 103.235.46.39:

```
    Packets: Sent = 4, Received = 4, Lost = 0 (0% loss),
Approximate round trip times in milli-seconds:
    Minimum = 3ms, Maximum = 3ms, Average = 3ms
```

C:\Users\rfan>nslookup

Default Server: intudns-vip.ied.edu.hk

Address: 192.168.166.103

> www.baidu.com

Server: intudns-vip.ied.edu.hk

Address: 192.168.166.103

Non-authoritative answer:

Name: www.wshifen.com

Address: 103.235.46.39

Aliases: www.baidu.com

www.a.shifen.com

# DNS: Domain Name System

*people:* many identifiers:

- SSN, name, passport #

*Internet hosts, routers:*

- IP address (32 bit) - used for addressing datagrams
- "name", e.g., cs.umass.edu - used by humans

Q: how to map between IP address and name, and vice versa ?

*Domain Name System:*

- *distributed database*  
implemented in hierarchy of many *name servers*
- *application-layer protocol:*  
hosts, name servers communicate to *resolve* names (address/name translation)
  - note: core Internet function, *implemented as application-layer protocol*
  - complexity at network's "edge"

# DNS: services, structure

## DNS services

- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - replicated Web servers: many IP addresses correspond to one name


## *Q: Why not centralize DNS?*

- single point of failure
- traffic volume
- distant centralized database
- maintenance

## *A: doesn't scale!*

- Comcast DNS servers alone: 600B DNS queries per day

# DNS: a distributed, hierarchical database



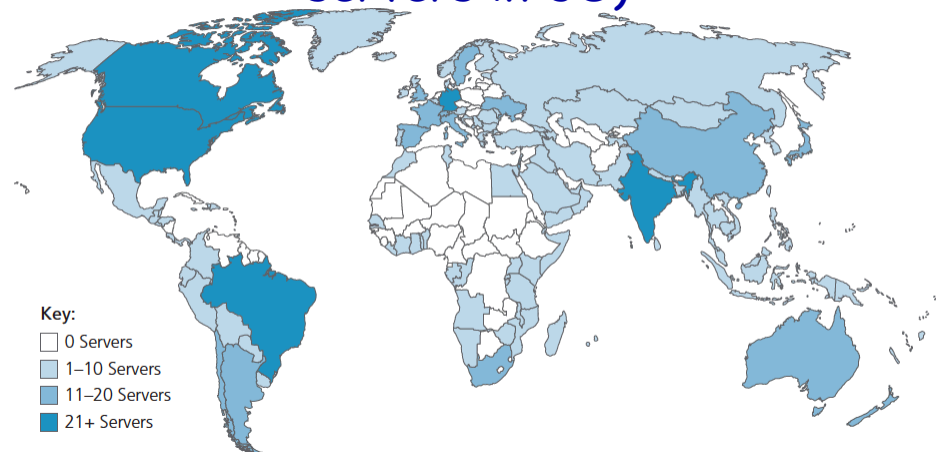
Client wants IP address for `www.amazon.com`; 1<sup>st</sup> approximation:

- client queries root server to find `.com` DNS server
- client queries `.com` DNS server to get `amazon.com` DNS server
- client queries `amazon.com` DNS server to get IP address for `www.amazon.com`

# DNS: root Name Servers

- official, contact-of-last-resort by name servers that can not resolve name
- *incredibly important* Internet function
  - Internet couldn't function without it!
  - DNSSEC - provides security (authentication and message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name "servers"  
worldwide each "server"  
replicated many times (~200  
servers in US)





# TLD: authoritative Servers

## Top-Level Domain (TLD) servers:

- responsible for .com, .org, .net, .edu, .aero, .jobs, .museums, and all top-level country domains, e.g.: .cn, .uk, .fr, .ca, .jp
- Network Solutions: authoritative registry for .com, .net TLD
- Educause: .edu TLD

## Authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

## Local DNS name Servers

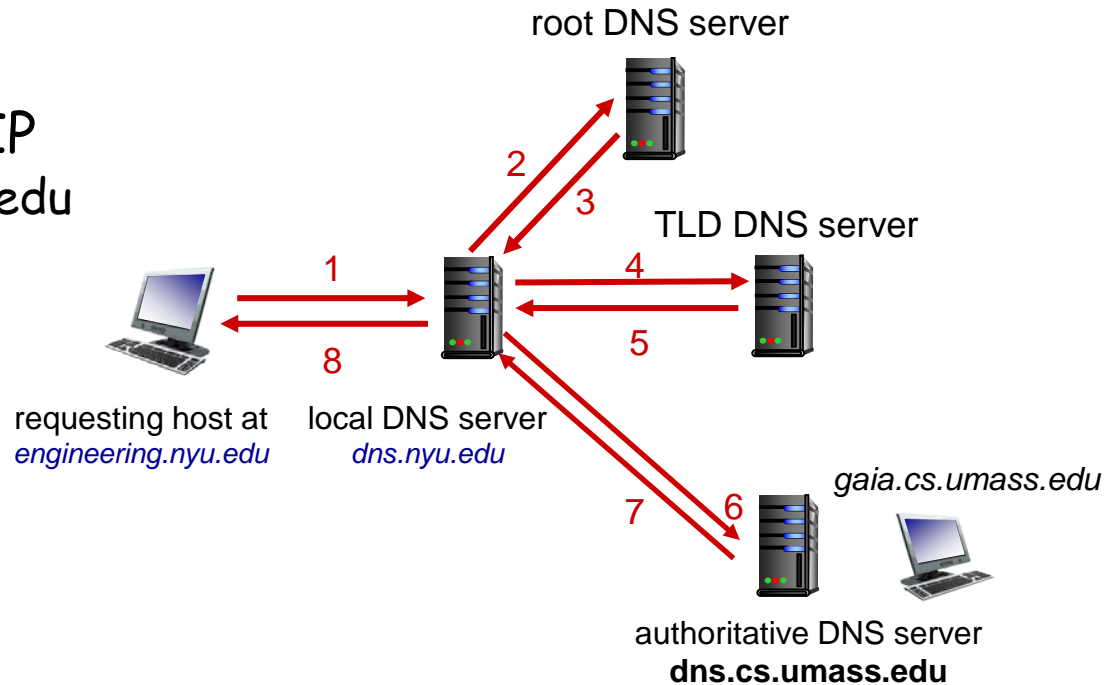
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
  - also called "default name server"
- when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date!)
  - acts as proxy, forwards query into hierarchy

# DNS name resolution: iterated query

**Example:** host at `engineering.nyu.edu` wants IP address for `gaia.cs.umass.edu`

## Iterated query:

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"

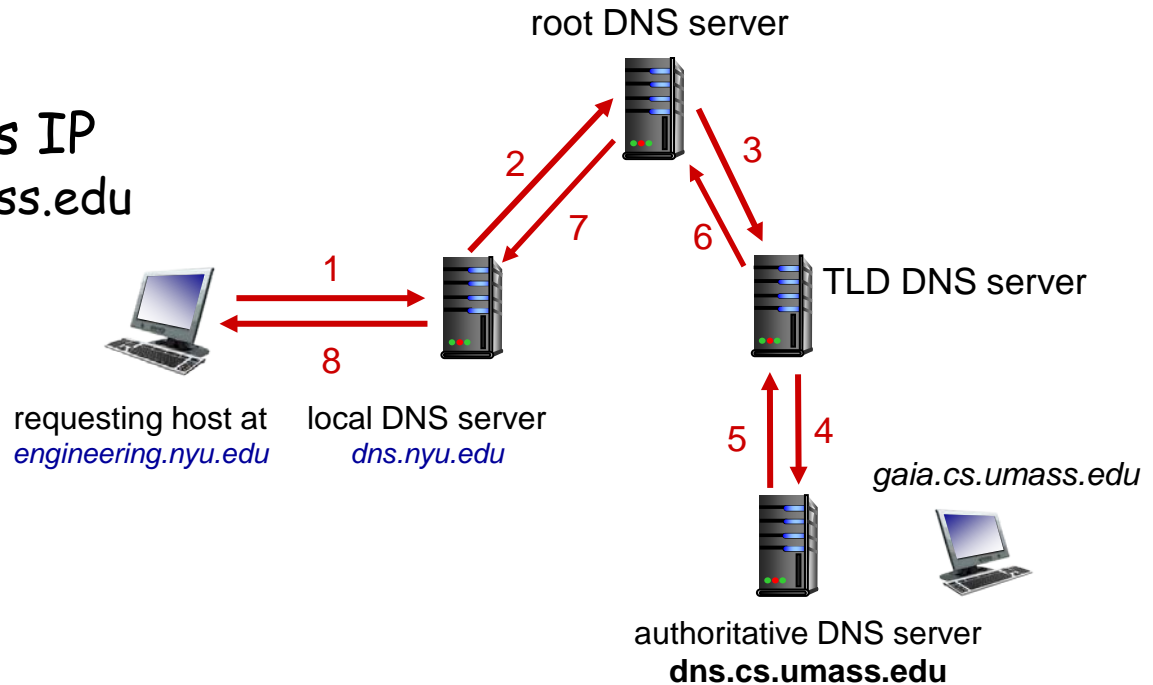


# DNS name resolution: recursive query

**Example:** host at `engineering.nyu.edu` wants IP address for `gaia.cs.umass.edu`

## Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



# Caching, Updating DNS Records

- once (any) name server learns mapping, it  *caches*  mapping
  - cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - thus root name servers not often visited
- cached entries may be  *out-of-date*  ( *best-effort name-to-address translation!* )
  - if name host changes IP address, may not be known Internet-wide until all TTLs expire!
- update/notify mechanisms proposed IETF standard
  - RFC 2136

# DNS security

## DDoS attacks

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

## Redirect attacks

- man-in-middle
  - intercept DNS queries
- DNS poisoning
  - send bogus replies to DNS server, which caches

DNSSEC  
[RFC 4033]

## Exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification



# Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System  
DNS
- P2P applications
- video streaming and content distribution networks
- **socket programming with UDP and TCP**

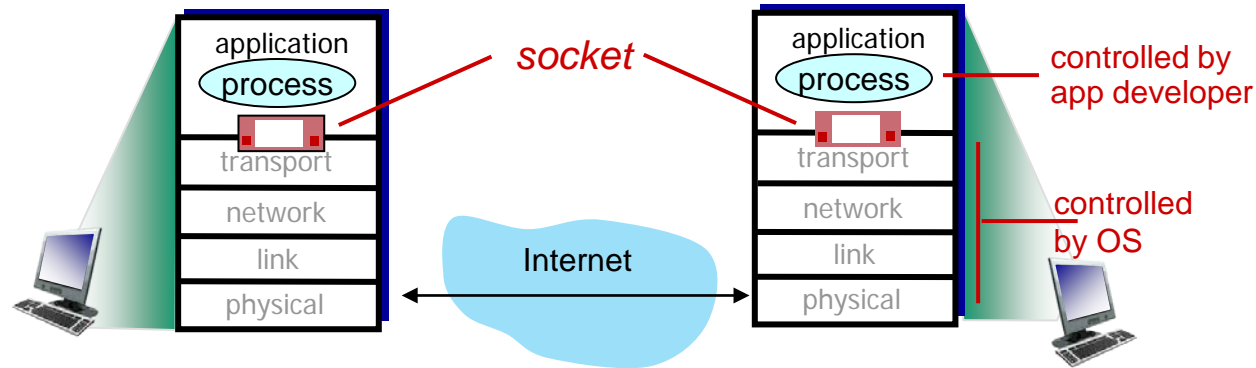




# Socket programming

*goal:* learn how to build client/server applications that communicate using sockets

*socket:* door between application process and end-end-transport protocol



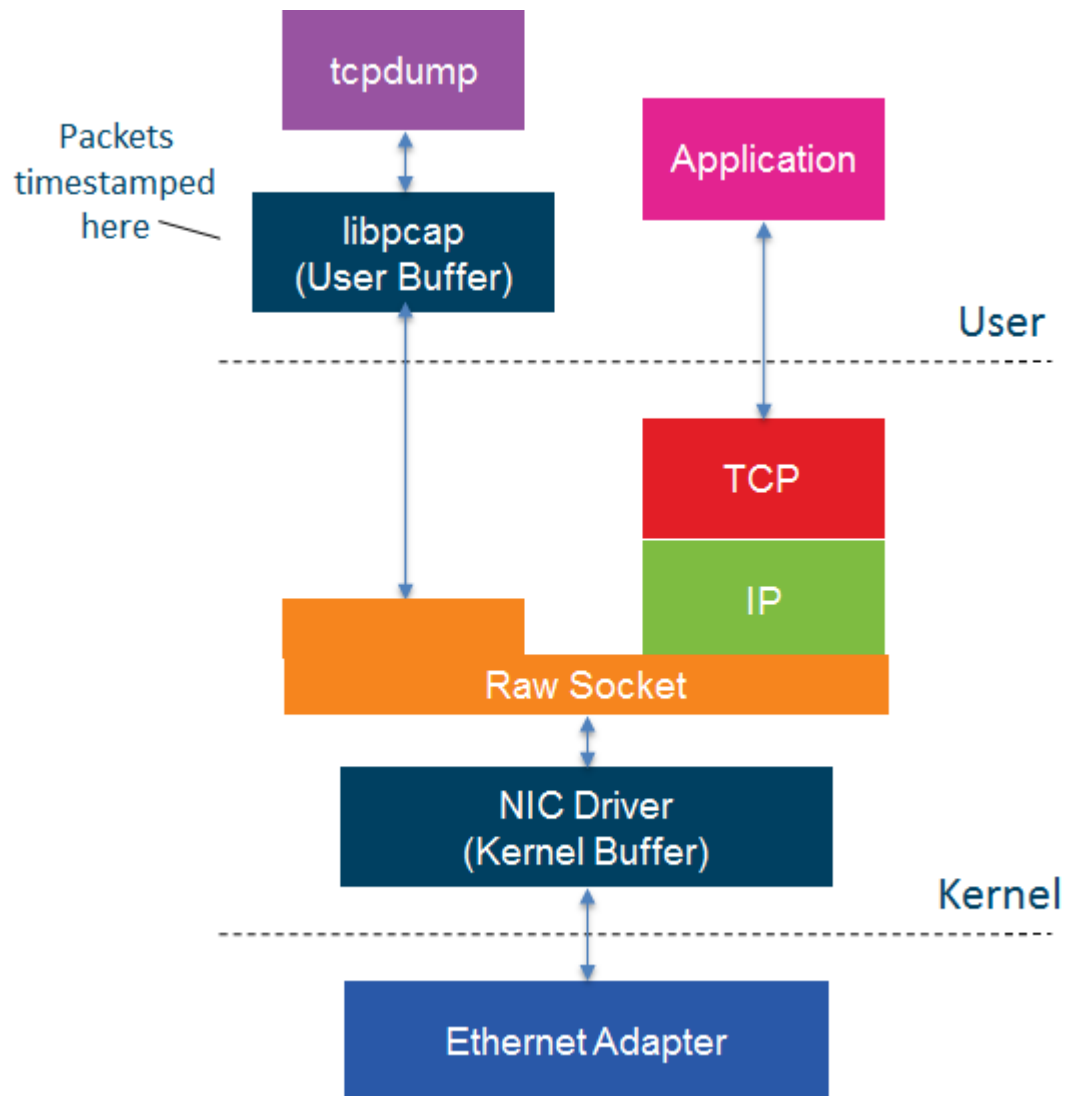
# Socket programming

Two socket types for two transport services:

- **UDP:** unreliable datagram
- **TCP:** reliable, byte stream-oriented

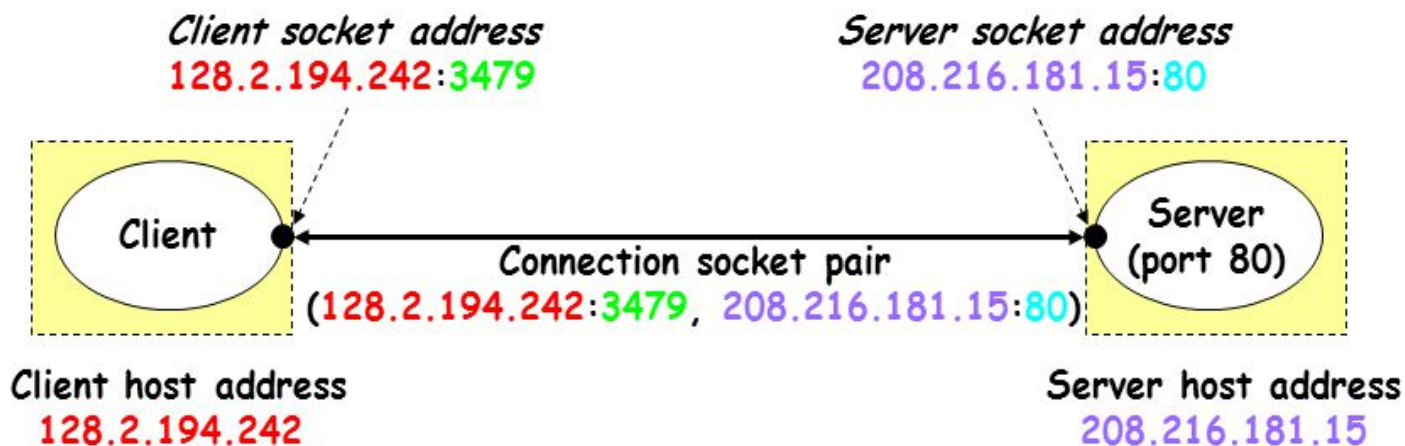
## Application Example:

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen



# Internet Connections (TCP/IP)

- Address the machine on the network
  - By IP address
- Address the process
  - By the "port"-number
- The pair of *IP-address + port* - makes up a "*socket-address*"



*Note: 3479 is an ephemeral port allocated by the kernel*

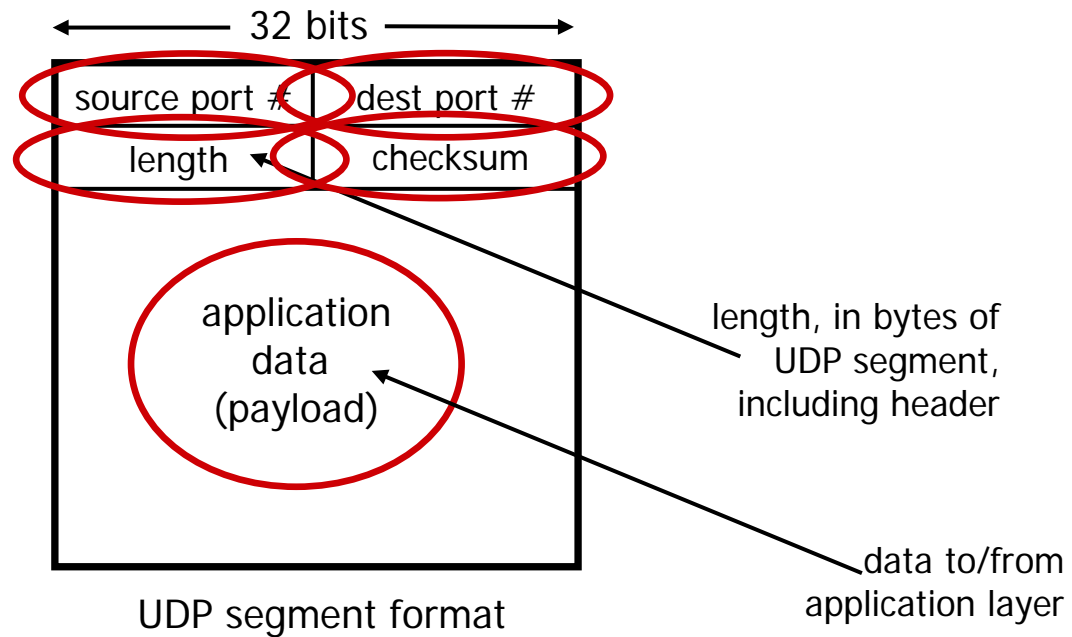
*Note: 80 is a well-known port associated with Web servers*

# What APIs Needed?

## Connection-Oriented TCP    Connectionless UDP

- How to create socket (door)
  - How to establish connection
    - Client connects to a server
    - Server accepts client req.
  - How to send/recv data
  - How to identify socket
    - Bind to local address/port
  - How to close socket (door)
- How to create socket
  - How to send/recv data
  - How to identify socket
  - How to close socket

# UDP segment header



# Socket programming with UDP

UDP: no “connection” between client & server

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

UDP: transmitted data may be lost or received out-of-order

Application viewpoint:

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

# Client/server socket interaction: UDP



**server** (running on serverIP)

create socket, port= x:  
`serverSocket =  
socket(AF_INET,SOCK_DGRAM)`

read datagram from  
`serverSocket`

write reply to  
`serverSocket`  
specifying  
client address,  
port number



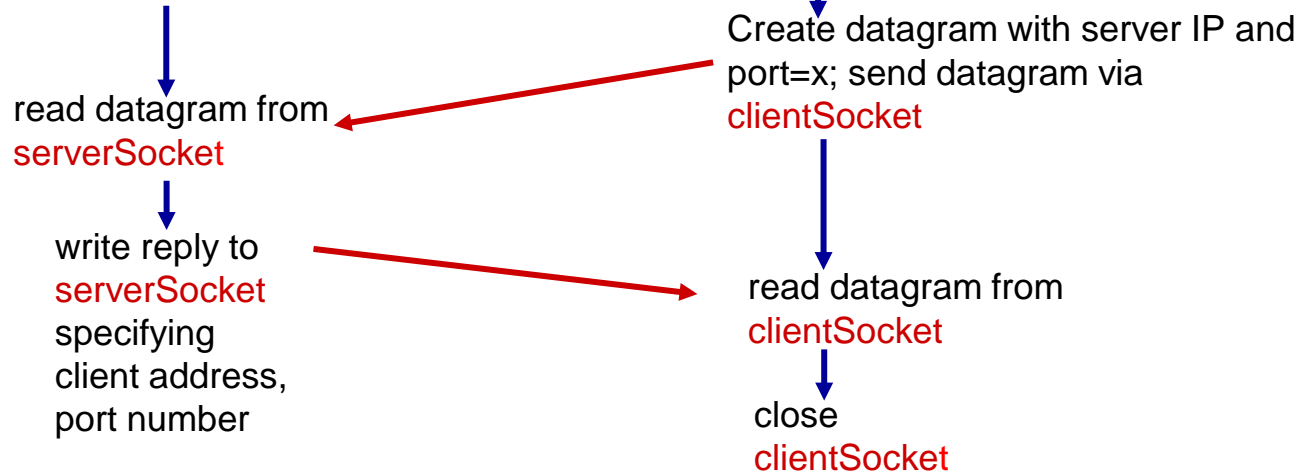
**client**

create socket:  
`clientSocket =  
socket(AF_INET,SOCK_DGRAM)`

Create datagram with server IP and  
port=x; send datagram via  
`clientSocket`

read datagram from  
`clientSocket`

close  
`clientSocket`





# Example app: UDP client

## *Python UDPClient*

include Python's socket library	→	from socket import *
		serverName = 'hostname'
		serverPort = 12000
create UDP socket for server	→	clientSocket = socket(AF_INET, SOCK_DGRAM)
get user keyboard input	→	message = raw_input('Input lowercase sentence:')
attach server name, port to message; send into socket	→	clientSocket.sendto(message.encode(), (serverName, serverPort))
read reply characters from socket into string	→	modifiedMessage, serverAddress = clientSocket.recvfrom(2048)
print out received string and close socket	→	print modifiedMessage.decode() clientSocket.close()

## Example app: UDP server

## Python UDPServer

	from socket import *
	serverPort = 12000
create UDP socket →	serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port number 12000 →	serverSocket.bind(('', serverPort))
	print ("The server is ready to receive")
loop forever →	while True:
Read from UDP socket into message, getting client's address (client IP and port) →	message, clientAddress = serverSocket.recvfrom(2048)
	modifiedMessage = message.decode().upper()
send upper case string back to this client →	serverSocket.sendto(modifiedMessage.encode(), clientAddress)

# TCP/IP Socket

## ■ Creating

- `int socket(int protocolFamily, int type, int protocol)`
  - protocol Family: `PF_INET`
  - type: `SOCK_STREAM`, `SOCK_DGRAM`
  - protocol: `IPPROTO_TCP`, `IPPROTO_UDP`
  - return value: File descriptor (socket descriptor)

## ■ Destroying

- `int close(int socket)`
  - socket: socket descriptor

# TCP/IP Socket

## ■ TCP Client

- Create a TCP socket
- Establish connection
  - `int connect(int socket, struct sockaddr *foreignAddress, unsigned int addressLength)`
- Communicate
  - `int send(int socket, const void *msg, unsigned int msgLength, int flags)`
  - `int recv(int socket, void *rcvBuffer, unsigned int bufferLength, int flags)`
- Close the connection

# TCP/IP Socket

## ■ TCP Server

- Create a TCP socket
- Assign a port to socket
  - `int bind(int socket, struct sockaddr *localAddress, unsigned int addressLength)`
- Set socket to listen
  - `int listen(int socket, int queueLimit)`
- Repeatedly
  - Accept new connection
    - `int accept(int socket, struct sockaddr *clientAddress, unsigned int *addressLength)`
  - Communicate
  - Close the connection

# Socket programming with TCP

## Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

## Client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket*: client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients (more in Chap 3)

## Application viewpoint

TCP provides reliable, in-order byte-stream transfer ("pipe") between client and server

# Client/server socket interaction: TCP



**server** (running on `hostid`)

**client**



create socket,  
port=`x`, for incoming  
request:  
`serverSocket = socket()`

wait for incoming  
connection request  
`connectionSocket =`  
`serverSocket.accept()`

read request from  
`connectionSocket`

write reply to  
`connectionSocket`

close  
`connectionSocket`

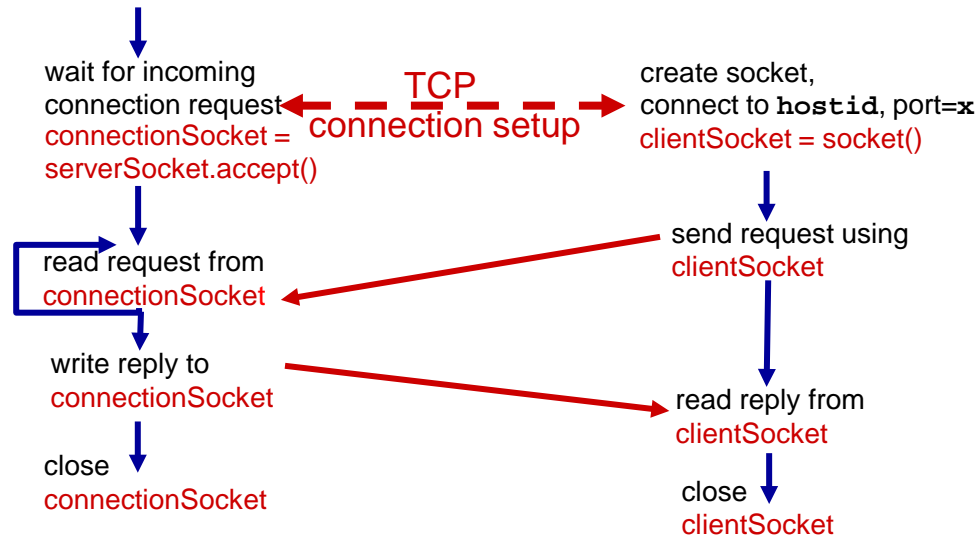
**TCP**  
connection setup

create socket,  
connect to `hostid`, port=`x`  
`clientSocket = socket()`

send request using  
`clientSocket`

read reply from  
`clientSocket`

close  
`clientSocket`



# Example app: TCP client

## *Python TCPClient*

create TCP socket for server,  
remote port 12000

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

No need to attach server name, port



# Example app: TCP server

## *Python TCPServer*

		<pre>from socket import *</pre>
		<pre>serverPort = 12000</pre>
create TCP welcoming socket	→	<pre>serverSocket = socket(AF_INET, SOCK_STREAM)</pre>
		<pre>serverSocket.bind(('', serverPort))</pre>
server begins listening for incoming TCP requests	→	<pre>serverSocket.listen(1)</pre>
		<pre>print 'The server is ready to receive'</pre>
		<pre>while True:</pre>
loop forever	→	<pre>connectionSocket, addr = serverSocket.accept()</pre>
server waits on accept() for incoming requests, new socket created on return	→	
		<pre>sentence = connectionSocket.recv(1024).decode()</pre>
read bytes from socket (but not address as in UDP)	→	<pre>capitalizedSentence = sentence.upper()</pre>
		<pre>connectionSocket.send(capitalizedSentence.encode())</pre>
close connection to this client (but <i>not</i> welcoming socket)	→	<pre>connectionSocket.close()</pre>

## Chapter 2: Summary

our study of network application layer is now complete!

- application architectures
  - client-server
  - P2P
- application service requirements:
  - reliability, bandwidth, delay
- Internet transport service model
  - connection-oriented, reliable: TCP
  - unreliable, datagrams: UDP
- specific protocols:
  - HTTP
  - SMTP, IMAP
  - DNS
  - P2P: BitTorrent
- video streaming, CDNs
- socket programming:  
TCP, UDP sockets

# Chapter 2: Summary

Most importantly: learned about *protocols*!

- typical request/reply message exchange:
  - client requests info or service
  - server responds with data, status code
- message formats:
  - *headers*: fields giving info about data
  - *data*: info(payload) being communicated

important themes:

- centralized vs. decentralized
- stateless vs. stateful
- scalability
- reliable vs. unreliable message transfer
- "complexity at network edge"

# Announcements

- 5/1 녹화동영상 강의 (Labor's Day – 학교 휴강 지침)
- 5/6 녹화동영상 강의 (대체공휴일)
- 5월 중 Quiz
- 5월 중하순 Midterm exam #2