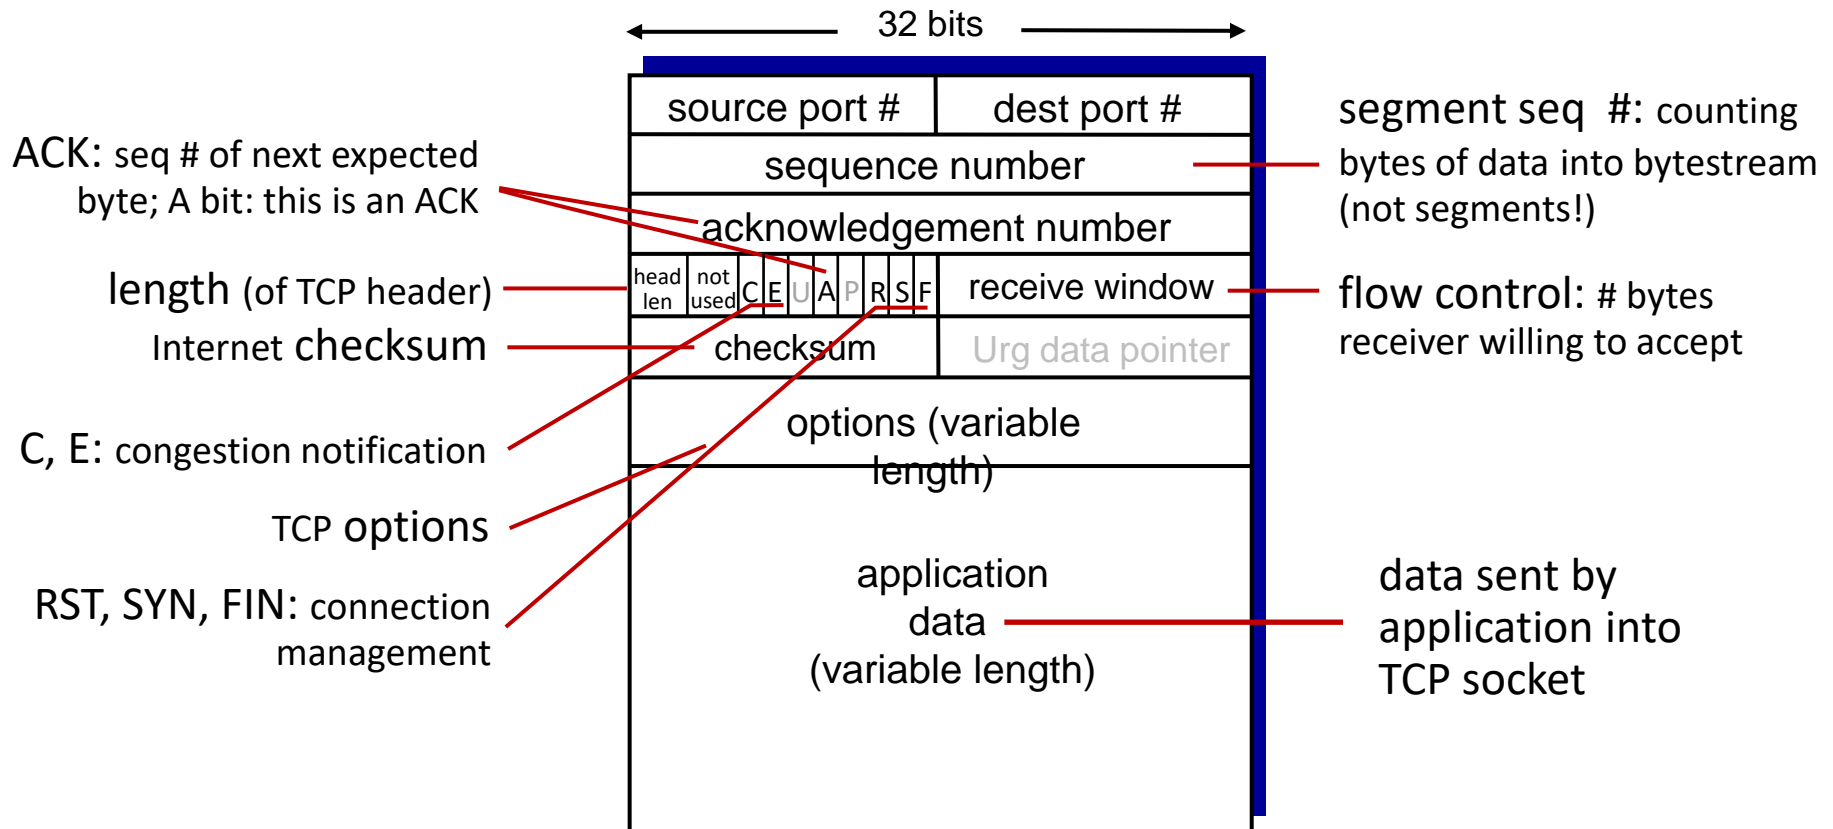


Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- **TCP congestion control**
- Evolution of transport-layer functionality



TCP segment structure



MSS

- Maximum Segment Size
- The default TCP Maximum Segment Size is 536
- For most computer users, the MSS option is established by the operating system
- The maximum segment size is specified as a TCP option, initially in the TCP SYN packet during the TCP handshake. The value cannot be changed after the connection is established

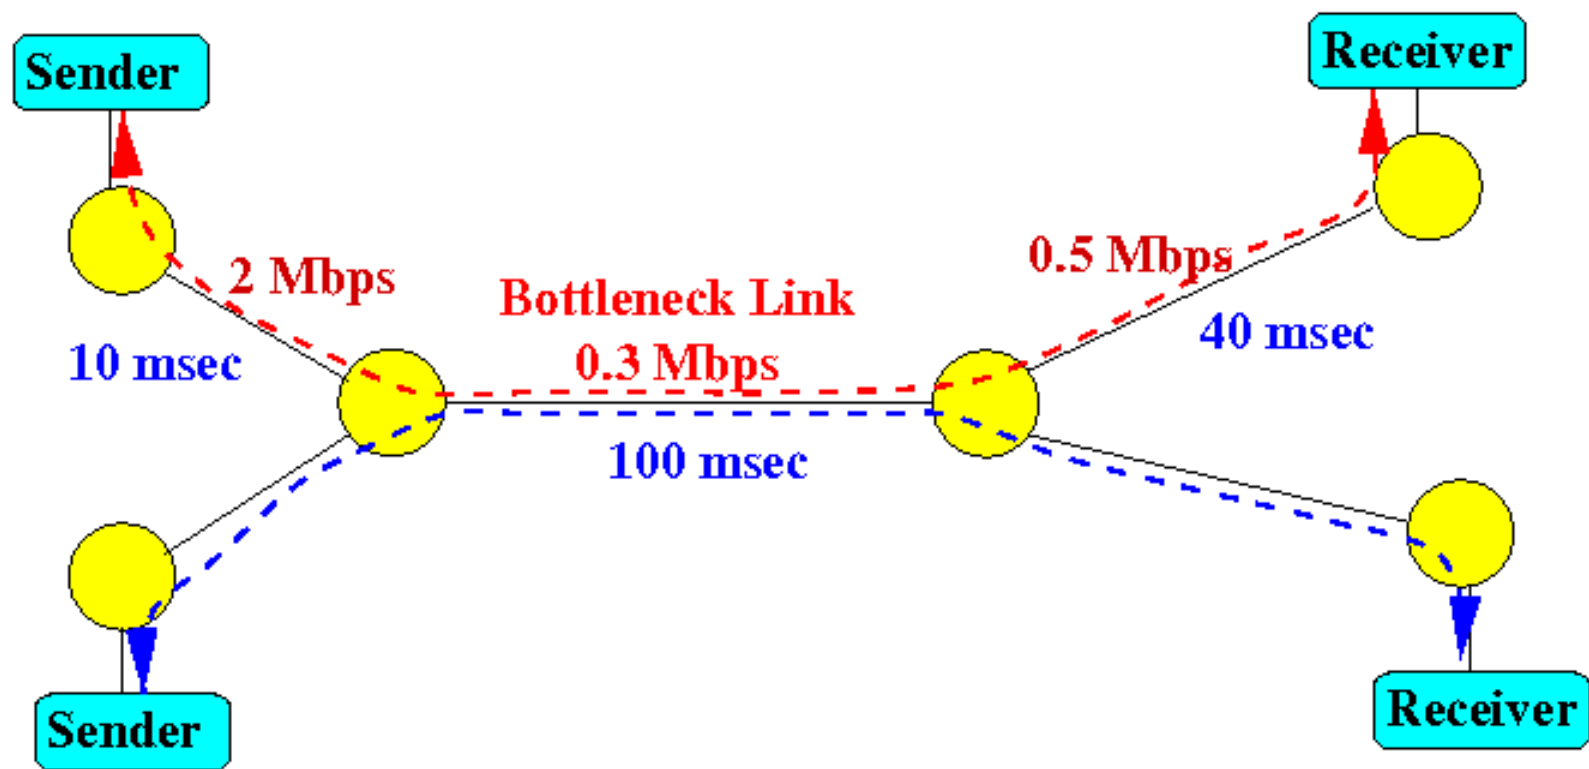
MSS


- To avoid fragmentation in the IP layer, a host must specify the maximum segment size as equal to the largest IP datagram that the host can handle minus the IP header size and TCP header sizes
- Therefore IPv4 hosts are required to be able to handle an MSS of 536 octets ($= 576 - 20 - 20$)

The TCP Maximum Segment Size and Related Topics (RFC 879, 1983)


The default IP Maximum Datagram Size is 576

The default TCP Maximum Segment Size is 536





TCP Behavior



RFC1323 - TCP Extensions for High Performance

TCP congestion control:

additive increase, multiplicative decrease

- *AIMD*

- *Approach:* increase transmission rate (***congestion window size***), probing for usable bandwidth, until loss occurs
 - *additive increase:* increase **cwnd** by 1 MSS every RTT until loss detected
 - *multiplicative decrease:* cut **cwnd** in half after loss

TCP congestion control: AIMD


- *approach*: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event

Additive Increase

increase sending rate by 1 maximum segment size every RTT until loss detected

Multiplicative Decrease

cut sending rate in half at each loss event



AIMD sawtooth behavior: *probing* for bandwidth

TCP AIMD: more

Multiplicative decrease detail: sending rate is


- Cut in half on loss detected by triple duplicate ACK (TCP Reno)
- Cut to 1 MSS (maximum segment size) when loss detected by timeout (TCP Tahoe)

Why AIMD?

- AIMD – a distributed, asynchronous algorithm – has been shown to:
 - optimize congested flow rates network wide!
 - have desirable stability properties

TCP slow start

- when connection begins, increase rate exponentially until first loss event:
 - initially **cwnd** = 1 MSS
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- *summary*: initial rate is slow, but ramps up exponentially fast



TCP: from slow start to congestion avoidance

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to 1/2 of its value before timeout.

Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event

