

# COSE321 Computer Systems Design

## Final Exam, Spring 2022

Name: Solutions

**Note: No Explanations, No Credits!**

1. You want to implement system calls (`fork()`, `open()`, and `close()`) via `svc` instruction and decided to use `'svc #1'` for `fork()`, `'svc #2'` for `open()` and `'svc #3'` for `close()`. Write the code in `SVC_ISR` that calls the appropriate handler. **(20 points)**

```
csd_vector_table:
```

```
    b .
    b SVC_ISR
    b .
    b .
    b .
    b .
    b .
    b .
```

```
// system call handlers
```

```
fork_handler:
```

```
    ...
    movs pc, lr
```

```
open_handler:
```

```
    ...
    movs pc, lr
```

```
close_handler:
```

```
    ...
    movs pc, lr
```

```
SVC_ISR:
```

```
sub    r1, lr, #4
ldr    r1, [r1] // read svc instruction
and    r1, r1, #0xFFFFFFFF
cmp    r1, #1
beq    fork_handler
cmp    r1, #2
beq    open_handler
cmp    r1, #3
beq    close_handler
```

### Encoding A1

ARMv4\*, ARMv5T\*, ARMv6\*, ARMv7

SVC<c> #<imm24>

|      |    |    |    |    |    |    |    |       |    |    |    |    |    |    |    |    |    |    |    |    |    |   |   |   |   |   |   |   |   |   |   |
|------|----|----|----|----|----|----|----|-------|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31   | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23    | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| cond |    |    |    | 1  | 1  | 1  | 1  | imm24 |    |    |    |    |    |    |    |    |    |    |    |    |    |   |   |   |   |   |   |   |   |   |   |

Table 1. LR adjustments

| Exceptions               | Adjustment | Returned Place       |
|--------------------------|------------|----------------------|
| Software Interrupt (SVC) | 0          | Next Instruction     |
| Undefined Instruction    | 0          | Next Instruction     |
| Prefetch Abort           | -4         | Aborting Instruction |
| Data Abort               | -8         | Aborting Instruction |
| FIQ                      | -4         | Next Instruction     |
| IRQ                      | -4         | Next Instruction     |

| Opcode<br>[31:28] | Mnemonic<br>extension | Meaning                           | Condition flag state  |
|-------------------|-----------------------|-----------------------------------|---|
| 0000              | EQ                    | Equal                             | Z set   |
| 0001              | NE                    | Not equal                         | Z clear   |
| 0010              | CS/HS                 | Carry set/unsigned higher or same | C set   |
| 0011              | CC/LO                 | Carry clear/unsigned lower        | C clear   |
| 0100              | MI                    | Minus/negative                    | N set   |
| 0101              | PL                    | Plus/positive or zero             | N clear   |
| 0110              | VS                    | Overflow                          | V set   |
| 0111              | VC                    | No overflow                       | V clear   |
| 1000              | HI                    | Unsigned higher                   | C set and Z clear   |
| 1001              | LS                    | Unsigned lower or same            | C clear or Z set  |
| 1010              | GE                    | Signed greater than or equal      | N set and V set, or<br>N clear and V clear (N == V)                             |
| 1011              | LT                    | Signed less than                  | N set and V clear, or<br>N clear and V set (N != V)                             |
| 1100              | GT                    | Signed greater than               | Z clear, and either N set and V set, or<br>N clear and V clear (Z == 0, N == V) |
| 1101              | LE                    | Signed less than or equal         | Z set, or N set and V clear, or<br>N clear and V set (Z == 1 or N != V)         |

2. Assume that there are two I/O devices (Private Timer and GPIO) generating normal interrupts (IRQ) to Arm CPU. Write the Arm assembly **1)** for setting up priority-related registers in GIC for the two I/Os and **2)** for jumping to an appropriate ISR (`Private_Timer_ISR` or `GPIO_ISR`) upon interrupt. The Private Timer's ID is 29, and the GPIO's ID is 61. Assume that the Private Timer has a higher priority than the GPIO. **Explain** your code. **(25 points)**

|   |  |
|---|--|
| <pre> // Interrupt Vector Table csd_vector_table:     b .     b .     b .     b .     b .     b .     b csd_IRQ_ISR     b .  main:  //Assume that VBAR and stack pointers //are set up appropriately beforehand  // 1) priority-related register //    set up in GIC //    (Distributor &amp; CPU Interface)      ldr r0, =GICD_PRIOR7     ldr r1, [r0] // 16 for ID# 29 (Private Timer)     mov r2, #0x10 &lt;&lt; 8     orr r1, r1, r2     str r1, [r0]      ldr r0, =GICD_PRIOR15     ldr r1, [r0] // 32 for ID# 61 (GPIO)     mov r2, #0x20 &lt;&lt; 8     orr r1, r1, r2     str r1, [r0]      // CPU Interface PMR     ldr r0, =GICC_PMR     mov r2, #0xFF      // Lowest     str r1, [r0] </pre> | <pre> // ----- // IRQ ISR // -----  csd_IRQ_ISR:  // 2) Jump to an appropriate ISR      // Read Interrupt Ack register     ldr r0, =GICC_IAR     ldr r3, [r0]      mov r4, #0x3FF     and r4, r3, r4     cmp r4, #29     beq Private_Timer_ISR     cmp r4, #61     beq GPIO_ISR     ...  // ISR for Private Timer Private_Timer_ISR:     ...     subs pc, lr, #4  // ISR for GPIO GPIO_ISR:     ...     subs pc, lr, #4 </pre> |
|---|--|

## <GIC Distributor>

**Base address: 0xF8F0\_1000**

**Table 4-1 Distributor register map**

| Offset      | Name  | Type            | Reset <sup>a</sup>                  | Description                                     |
|-------------|---|-----------------|-------------------------------------|---|
| 0x000       | <a href="#">GICD_CTLR</a>                   | RW              | 0x00000000                          | Distributor Control Register                    |
| 0x004       | <a href="#">GICD_TYPER</a>                  | RO              | IMPLEMENTATION DEFINED              | Interrupt Controller Type Register              |
| 0x008       | <a href="#">GICD_IIDR</a>                   | RO              | IMPLEMENTATION DEFINED              | Distributor Implementer Identification Register |
| 0x00C-0x01C | -   | -               | -                                   | Reserved  |
| 0x020-0x03C | -   | -               | -                                   | IMPLEMENTATION DEFINED registers                |
| 0x040-0x07C | -   | -               | -                                   | Reserved  |
| 0x080       | <a href="#">GICD_IGROUPRn<sup>b</sup></a>   | RW              | IMPLEMENTATION DEFINED <sup>c</sup> | Interrupt Group Registers                       |
| 0x084-0x0FC |   |                 | 0x00000000                          |   |
| 0x100-0x17C | <a href="#">GICD_ISENABLERn</a>             | RW              | IMPLEMENTATION DEFINED              | Interrupt Set-Enable Registers                  |
| 0x180-0x1FC | <a href="#">GICD_ICENABLERn</a>             | RW              | IMPLEMENTATION DEFINED              | Interrupt Clear-Enable Registers                |
| 0x200-0x27C | <a href="#">GICD_ISPENDRn</a>               | RW              | 0x00000000                          | Interrupt Set-Pending Registers                 |
| 0x280-0x2FC | <a href="#">GICD_ICPENDRn</a>               | RW              | 0x00000000                          | Interrupt Clear-Pending Registers               |
| 0x300-0x37C | <a href="#">GICD_ISACTIVERn<sup>d</sup></a> | RW              | 0x00000000                          | GICv2 Interrupt Set-Active Registers            |
| 0x380-0x3FC | <a href="#">GICD_ICACTIVERn<sup>e</sup></a> | RW              | 0x00000000                          | Interrupt Clear-Active Registers                |
| 0x400-0x7F8 | <a href="#">GICD_IPRIORITYRn</a>            | RW              | 0x00000000                          | Interrupt Priority Registers                    |
| 0x7FC       | -   | -               | -                                   | Reserved  |
| 0x800-0x81C | <a href="#">GICD_ITARGETSRn</a>             | RO <sup>f</sup> | IMPLEMENTATION DEFINED              | Interrupt Processor Targets Registers           |
| 0x820-0xBF8 |   | RW <sup>f</sup> | 0x00000000                          |   |
| 0x8FC       | -   | -               | -                                   | Reserved  |
| 0xC00-0xCFC | <a href="#">GICD_ICFGRn</a>                 | RW              | IMPLEMENTATION DEFINED              | Interrupt Configuration Registers               |
| 0xD00-0xDFC | -   | -               | -                                   | IMPLEMENTATION DEFINED registers                |
| 0xE00-0xEFC | <a href="#">GICD_NSACRn<sup>e</sup></a>     | RW              | 0x00000000                          | Non-secure Access Control Registers, optional   |
| 0xF00       | <a href="#">GICD_SGIR</a>                   | WO              | -                                   | Software Generated Interrupt Register           |
| 0xF04-0xF0C | -   | -               | -                                   | Reserved  |
| 0xF10-0xF1C | <a href="#">GICD_CPENDSGIRn<sup>e</sup></a> | RW              | 0x00000000                          | SGI Clear-Pending Registers                     |
| 0xF20-0xF2C | <a href="#">GICD_SPENDSGIRn<sup>e</sup></a> | RW              | 0x00000000                          | SGI Set-Pending Registers                       |
| 0xF30-0xFCC | -   | -               | -                                   | Reserved  |
| 0xFD0-0xFFC | -   | RO              | IMPLEMENTATION DEFINED              | <i>Identification registers on page 4-119</i>   |

- a. For details of any restrictions that apply to the reset values of IMPLEMENTATION DEFINED cases see the appropriate register description.
- b. In a GICv1 implementation, present only if the GIC implements the GIC Security Extensions, otherwise RAZ/WI.
- c. For more information see [GICD\\_IGROUPR0 reset value on page 4-92](#).
- d. In GICv1, these are the Active Bit Registers, ICDABRn. These registers are RO.

## <GIC CPU Interface>

### Table 4-2 CPU interface register map

| Offset        | Name   | Type | Reset                   | Description   |
|---------------|--|------|-------------------------|---|
| 0x0000        | <a href="#">GICC_CTLR</a>                            | RW   | 0x00000000              | CPU Interface Control Register                      |
| 0x0004        | <a href="#">GICC_PMR</a>                             | RW   | 0x00000000              | Interrupt Priority Mask Register                    |
| 0x0008        | <a href="#">GICC_BPR</a>                             | RW   | 0x0000000x <sup>a</sup> | Binary Point Register                               |
| 0x000C        | <a href="#">GICC_IAR</a>                             | RO   | 0x000003FF              | Interrupt Acknowledge Register                      |
| 0x0010        | <a href="#">GICC_EOIR</a>                            | WO   | -                       | End of Interrupt Register                           |
| 0x0014        | <a href="#">GICC_RPR</a>                             | RO   | 0x000000FF              | Running Priority Register                           |
| 0x0018        | <a href="#">GICC_HPPIR</a>                           | RO   | 0x000003FF              | Highest Priority Pending Interrupt Register         |
| 0x001C        | <a href="#">GICC_ABPR</a> <sup>b</sup>               | RW   | 0x0000000x <sup>a</sup> | Aliased Binary Point Register                       |
| 0x0020        | <a href="#">GICC_AIAR</a> <sup>c</sup>               | RO   | 0x000003FF              | Aliased Interrupt Acknowledge Register              |
| 0x0024        | <a href="#">GICC_AEOIR</a> <sup>c</sup>              | WO   | -                       | Aliased End of Interrupt Register                   |
| 0x0028        | <a href="#">GICC_AHPPIR</a> <sup>c</sup>             | RO   | 0x000003FF              | Aliased Highest Priority Pending Interrupt Register |
| 0x002C–0x003C | -  | -    | -                       | Reserved  |
| 0x0040–0x00CF | -  | -    | -                       | IMPLEMENTATION DEFINED registers                    |
| 0x00D0–0x00DC | <a href="#">GICC_APR</a> <sup>n</sup> <sup>c</sup>   | RW   | 0x00000000              | Active Priorities Registers                         |
| 0x00E0–0x00EC | <a href="#">GICC_NSAPR</a> <sup>n</sup> <sup>c</sup> | RW   | 0x00000000              | Non-secure Active Priorities Registers              |
| 0x00ED–0x00F8 | -  | -    | -                       | Reserved  |
| 0x00FC        | <a href="#">GICC_IIDR</a>                            | RO   | IMPLEMENTATION DEFINED  | CPU Interface Identification Register               |
| 0x1000        | <a href="#">GICC_DIR</a> <sup>c</sup>                | WO   | -                       | Deactivate Interrupt Register                       |

a. See the register description for more information.

b. Present in GICv1 if the GIC implements the GIC Security Extensions. Always present in GICv2.

c. GICv2 only.

Figure 4-27 shows the IAR bit assignments.



**Figure 4-27 GICC\_IAR bit assignments**

Table 4-34 shows the IAR bit assignments.

#### Table 4-34 GICC\_IAR bit assignments

| Bit     | Name         | Function   |
|---------|--------------|--|
| [31:13] | -            | Reserved.  |
| [12:10] | CPUID        | For SGIs in a multiprocessor implementation, this field identifies the processor that requested the interrupt. It returns the number of the CPU interface that made the request, for example a value of 3 means the request was generated by a write to the <a href="#">GICD_SGIR</a> on CPU interface 3.<br>For all other interrupts this field is RAZ. |
| [9:0]   | Interrupt ID | The interrupt ID.  |

3. Refer to the following code. Assume that `VAR` is already set correctly for using our vector table, and all stack pointers are also set appropriately beforehand. **(30 points)**
- a. While CPU is executing `forever` loop, assume that a hardware interrupt is delivered to Arm CPU via the IRQ input. **Explain the execution flow** of the code from `main`. **(15 points)**

|   |   |
|---|---|
| <pre>csd_vector_table:     b .     b csd_undefined     b .     b .     b .     b csd_IRQ_ISR     b .  main:     cpsIE aif, #0x10 // Change to User Mode</pre> | <pre>forever:     b forever  csd_undefined:     movs pc, lr  csd_IRQ_ISR:      .word 0xffffffff     subs pc, lr, #4</pre> |
|---|---|

Execution flow:

main → csd\_IRQ\_ISR → Undefined exception → return to csd\_IRQ\_ISR → return to main

- b. While CPU is executing `forever` loop, assume that a hardware interrupt is delivered to Arm CPU via the IRQ input. Then another hardware interrupt is delivered to Arm CPU via the IRQ input while executing `nops` in ISR. **Explain the execution flow** of the code from `main`. **(15 points)**

|   |  |
|---|--|
| <pre>csd_vector_table:     b .     b .     b .     b .     b .     b .     b .     b csd_IRQ_ISR     b .  main:     cpsIE aif, #0x10 // Change to User Mode</pre> | <pre>forever:     b forever  csd_IRQ_ISR:      cpsIE i      nop     nop     nop      subs pc, lr, #4</pre> |
|---|--|

**Execution flow:**

`main` → `csd_IRQ_ISR` (1<sup>st</sup> interrupt) → `csd_IRQ_ISR` (2<sup>nd</sup> interrupt) → return to IRQ handler → return to IRQ handler (not `forever` loop in main code)

4. Figure out the address translation sizes and mappings from virtual address to physical address, with the following page tables. Refer to the page table entry information in the next page. Assume TTBR0 is set correctly beforehand. **Explain** your answer in detail. **(20 points)**

```
// 1st level page table (csd_MMUTable)

csd_MMUTable:
.set SECT, 0xAAA00000
.word SECT + 0x15de6
.word csd_MMUTable_lv2_1 + 0x1e1
.word csd_MMUTable_lv2_2 + 0x1e1

// 2nd level page tables (csd_MMUTable_lv2_1, csd_MMUTable_lv2_2)

csd_MMUTable_lv2_1:
.word 0xAAA00002

csd_MMUTable_lv2_2:
.word 0xAAA00002
```

| Translation size (4KB, 64KB, 1MB, or 16MB?) | Virtual address range     |   | Physical address range    |
|---|---------------------------|---|---------------------------|
| 1MB   | 0x0000_0000 ~ 0x000F_FFFF | → | 0xAAA0_0000 ~ 0xAAAF_FFFF |
| 4KB   | 0x0010_0000 ~ 0x0010_0FFF |   | 0xAAA0_0000 ~ 0xAAA0_0FFF |
| 4KB   | 0x0020_0000 ~ 0x0020_0FFF |   | 0xAAA0_0000 ~ 0xAAA0_0FFF |



## Short-descriptor translation table first-level descriptor formats

Each entry in the first-level table describes the mapping of the associated 1MB MVA range.

Figure B3-4 shows the possible first-level descriptor formats.

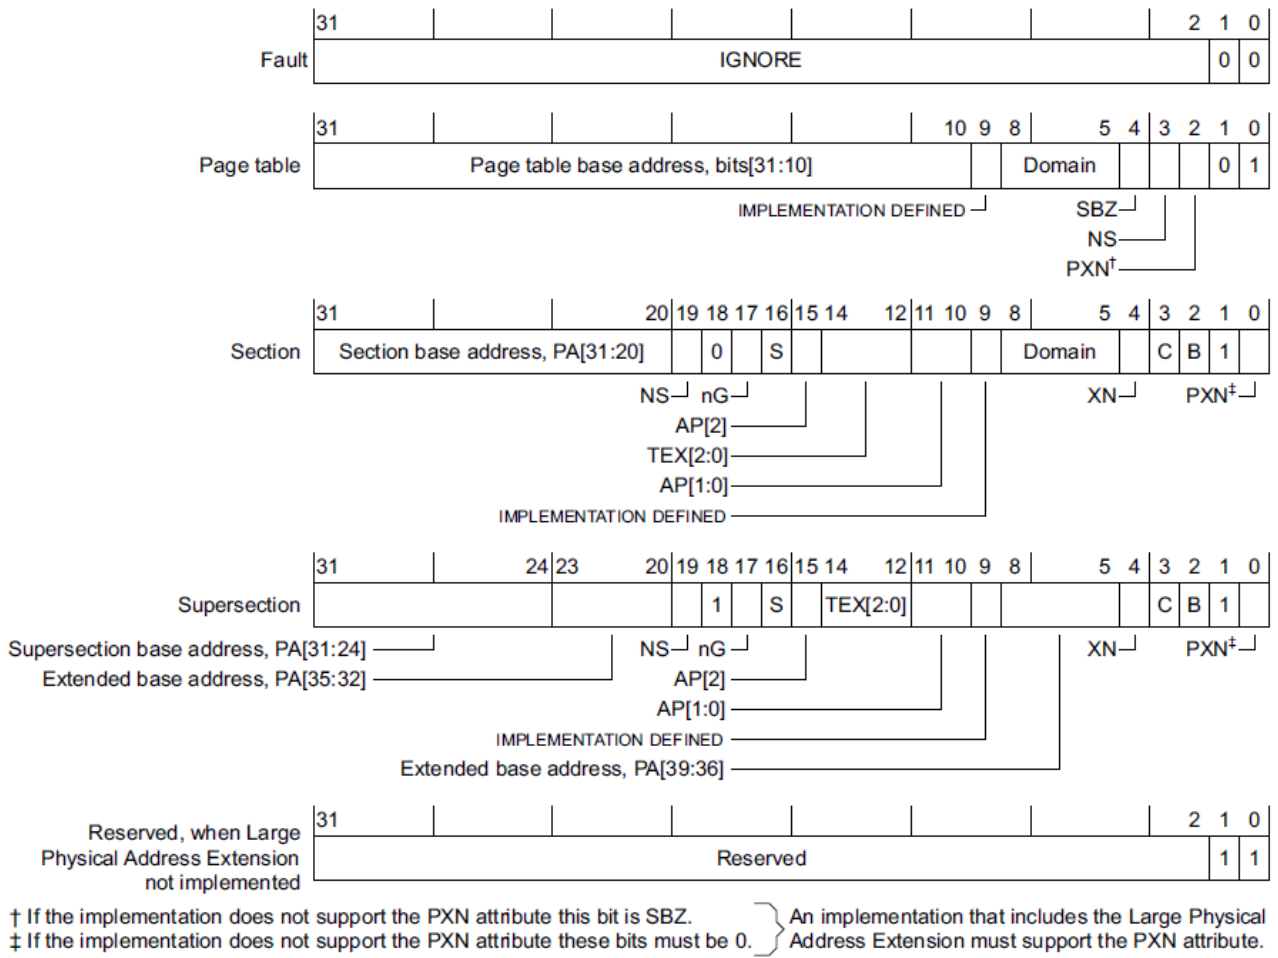


Figure B3-4 Short-descriptor first-level descriptor formats

## Short-descriptor translation table second-level descriptor formats

Figure B3-5 shows the possible formats of a second-level descriptor.

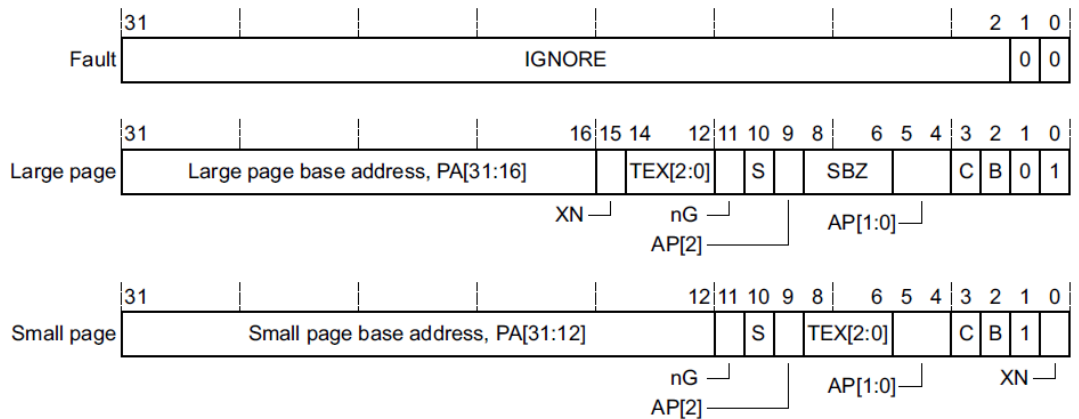


Figure B3-5 Short-descriptor second-level descriptor formats

5. Assume that the monitor node was entered via `smc` instruction. **(25 points)**

a. Explain what operations are performed by CPU (hardware) upon the execution of `smc` instruction. **(10 points)**

- $SPSR\_mon \leftarrow CPSR$
- $lr\_mon\ (r14) \leftarrow \text{address of next inst.}$
- $CPSR.mode \leftarrow \text{monitor mode (10110)}$
- $CPSR.A, I, F = 111$
- $PC \leftarrow MVBAR + 0x8$

b. You want to do both **world switching** and **context-switching** in the monitor mode. The context switching code is prepared for you below. Write code of performing the world-switching and calling an appropriate context-switching function. **(15 points)**

```
// Exception Vector Table in Monitor mode
csd_monitor_vector:
    b .
    b .
    b csd_SMC_handler
    b .
    b .
    b .
    b .
    b .

csd_SMC_Handler:

// Write the code of 1) performing the world-switching and
//                               2) calling an appropriate context-switching function

mrc    p15, 0, r0, c1, c1, 0 ; Read SCR
tst     r0, #1                ; Is the NS bit set?
eor     r0, r0, #b'1          ; Toggle NS bit
mcr     p15, 0, r0, c1, c1, 0 ; Write to SCR
beq     push_secure_pop_normal
bne     push_normal_pop_secure

// context-switching code below

// push the normal world context to stack
// pop the secure world context from stack
push_Normal_pop_Secure:
    ...

// push the secure world context to stack
// pop the normal world context from stack
push_Secure_pop_Normal:
    ...
```

### Accessing the SCR

To access the SCR, software reads or writes the CP15 registers with `<opc1>` set to 0, `<CRn>` set to c1, `<CRm>` set to c1, and `<opc2>` set to 0. For example:

```
MRC p15, 0, <Rt>, c1, c1, 0    ; Read SCR into Rt
MCR p15, 0, <Rt>, c1, c1, 0    ; Write Rt to SCR
```